

AthenaX: Streaming Processing Platform @Uber

Bill Liu, Haohui Mai

APRIL 11, 2017

The Uber logo, consisting of the word "UBER" in white, uppercase, sans-serif font, centered within a black rectangular box.

UBER




Speakers



Bill Liu

- Senior Software Engineer @ Uber

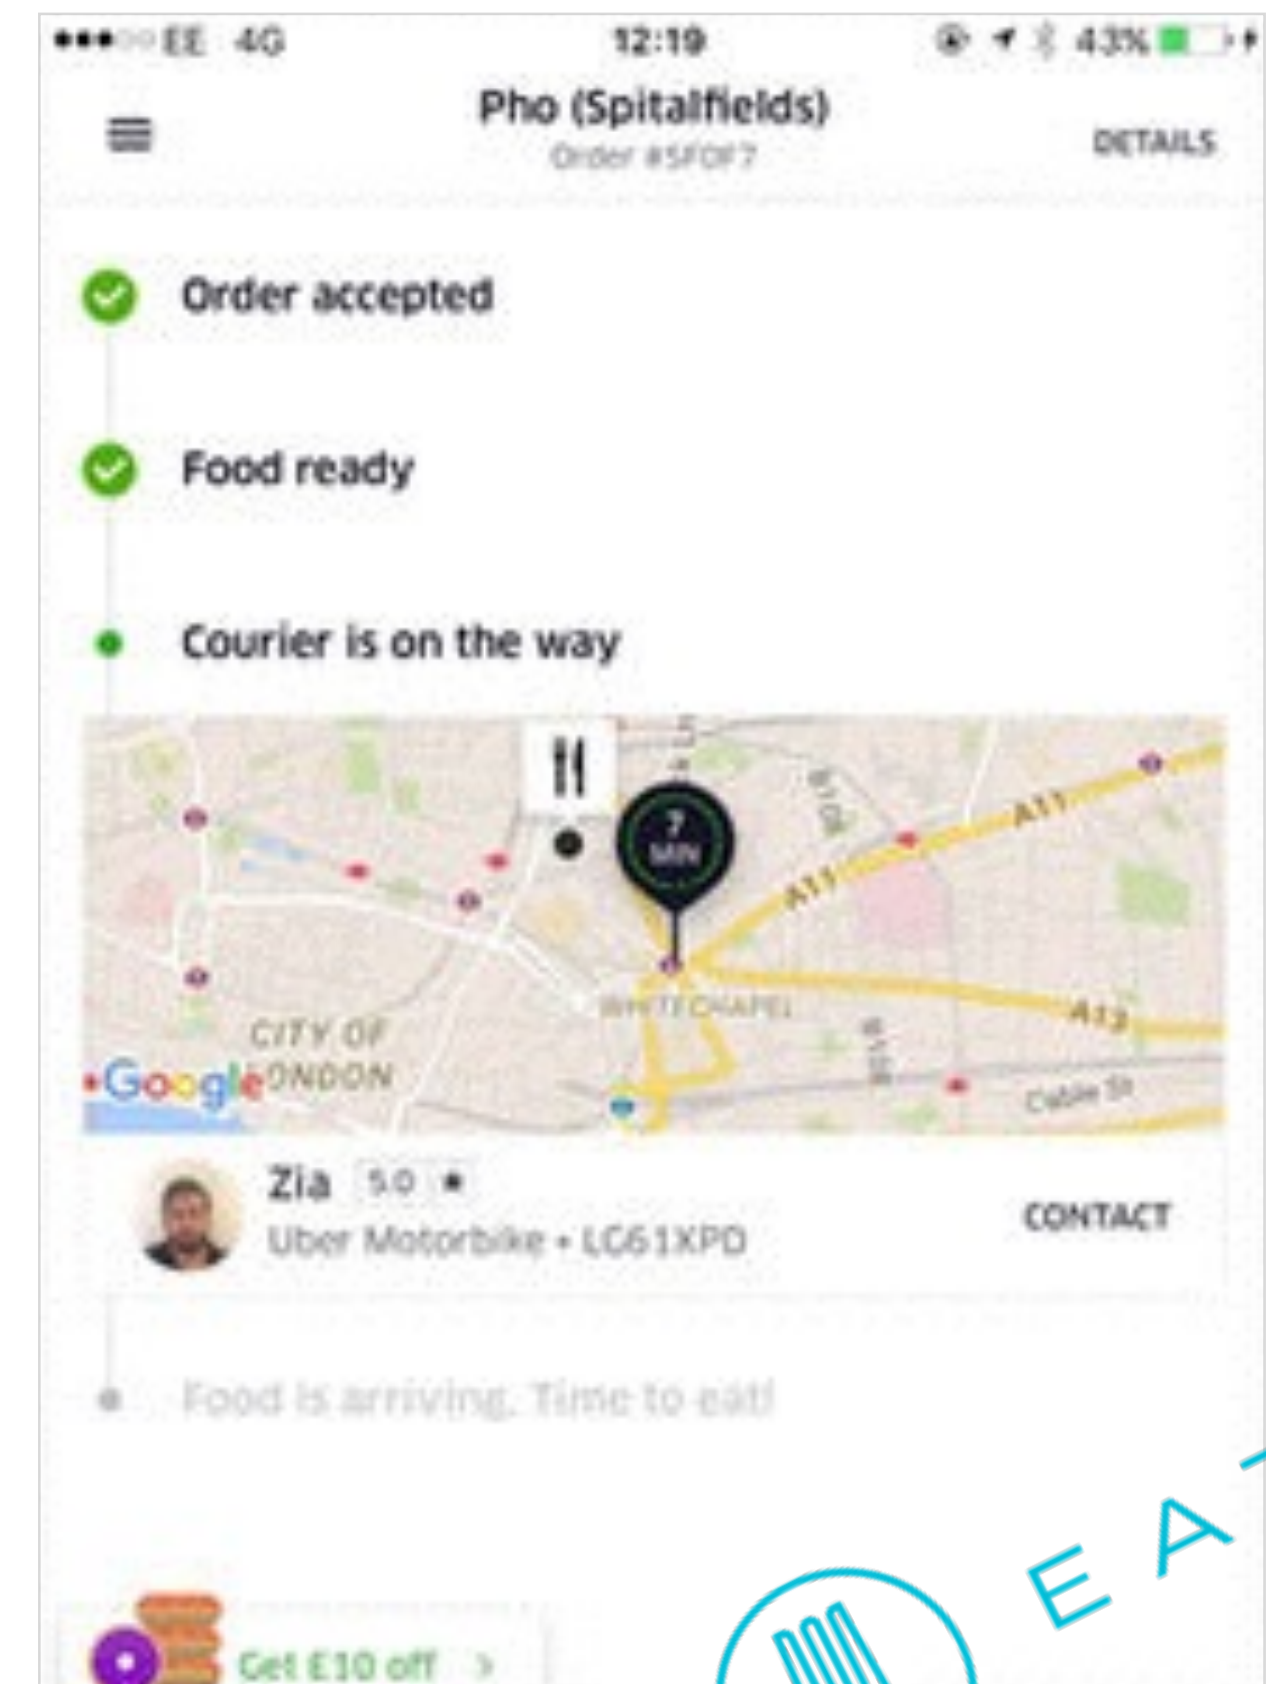


Haohui Mai, @wheat9

- Senior Software Engineer @ Uber
- PMC, Apache Hadoop & Storm

Uber business is real-time

- Uber: Transport A → B on demand reliably
- Dynamic marketplace
- Example: UberEATS



Challenges

Infra.: Reliability & scalability

- 99.99% SLA on latency
- At-least-once processing
- Billions of messages
- Multiple PB / day

Solutions: Productivity

- Audiences: majority of employees use SQL actively
- Abstractions: Flink / DSL?
- Integrations: data management, monitoring, reporting, etc.

Building streaming applications



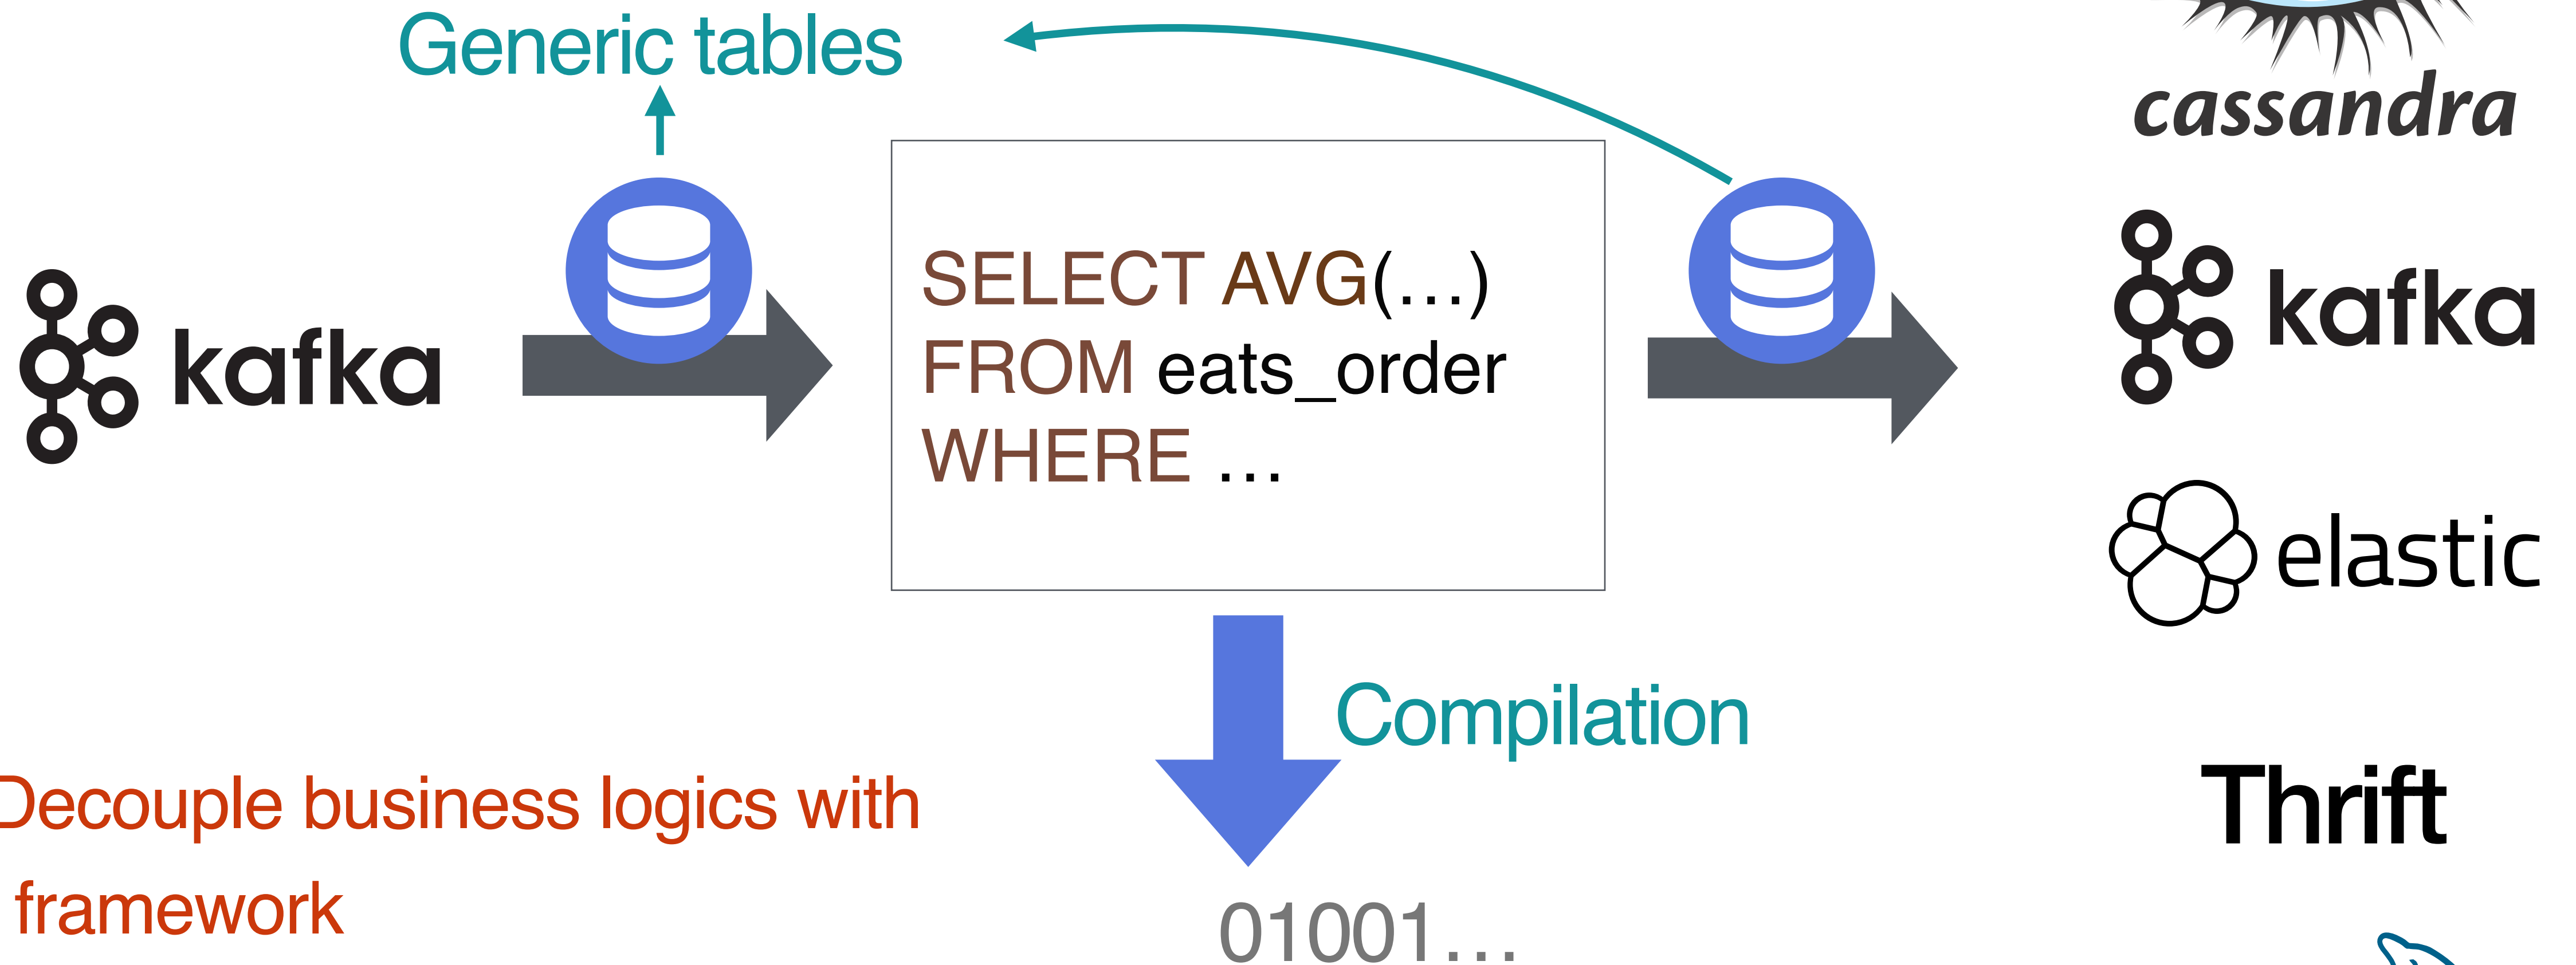
Thrift



- Framework-specific
- Ad-hoc management over the life-cycles

The AthenaX approach

Write SQLs to build streaming applications



- Decouple business logics with framework
- Unified integration & management

 **memsql**

 **cassandra**

 **kafka**

 **elastic**

Thrift

 **MySQL**

AthenaX: Streaming processing platform @ Uber

- Write SQLs to build streaming applications
 - **Insight: generic table**
- Reliable, scalable processing based on Apache Flink
- Develop & deploy streaming applications in production in **hours** instead of weeks

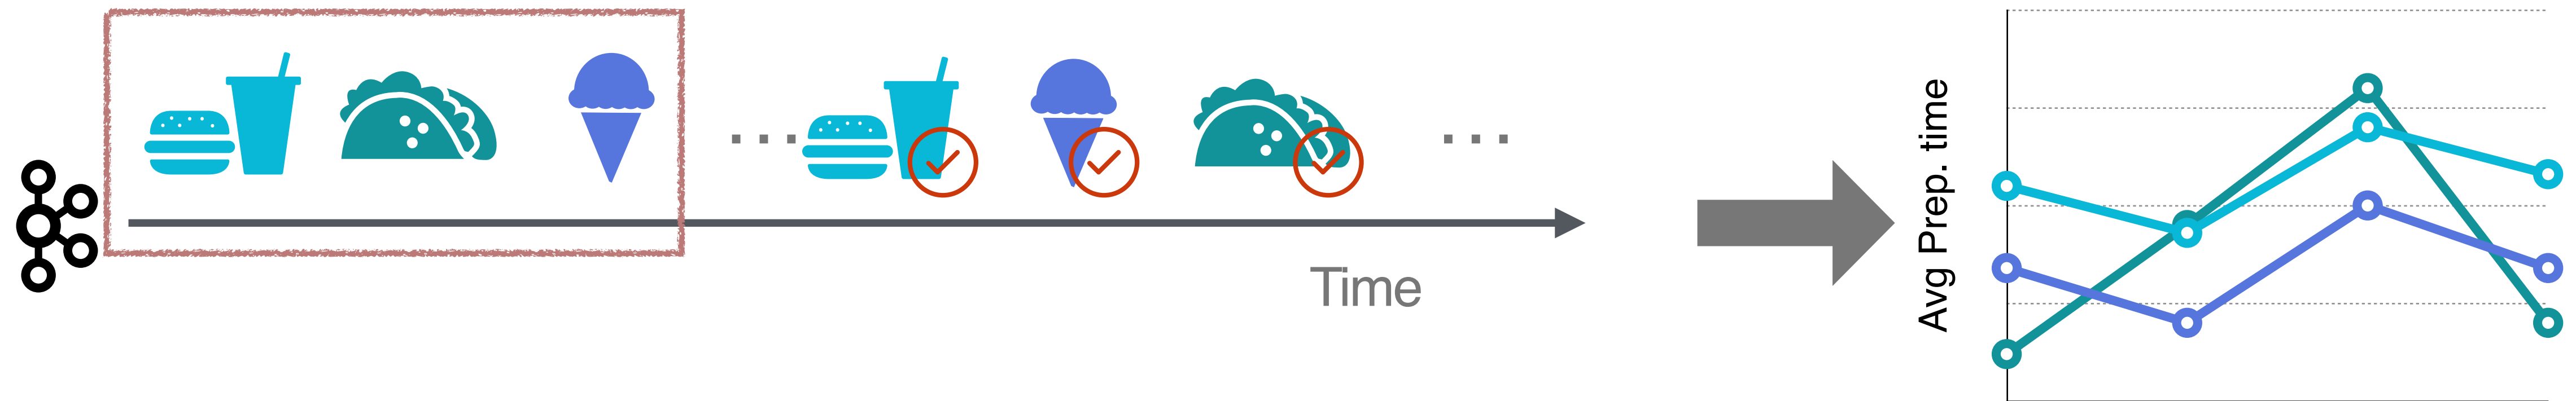


Agenda

- Motivating example
- Case study: ETD in UberEATS
- Implementation
- Current status
- Conclusion

Example

Real-time dashboard for restaurants



```
SELECT meal_id, AVG(meal_prep_time)
FROM eats_order
```

```
GROUP BY meal_id, HOP(proctime(),
  INTERVAL '1' MINUTE,
  INTERVAL '15' MINUTE)
```

Example (cont.)

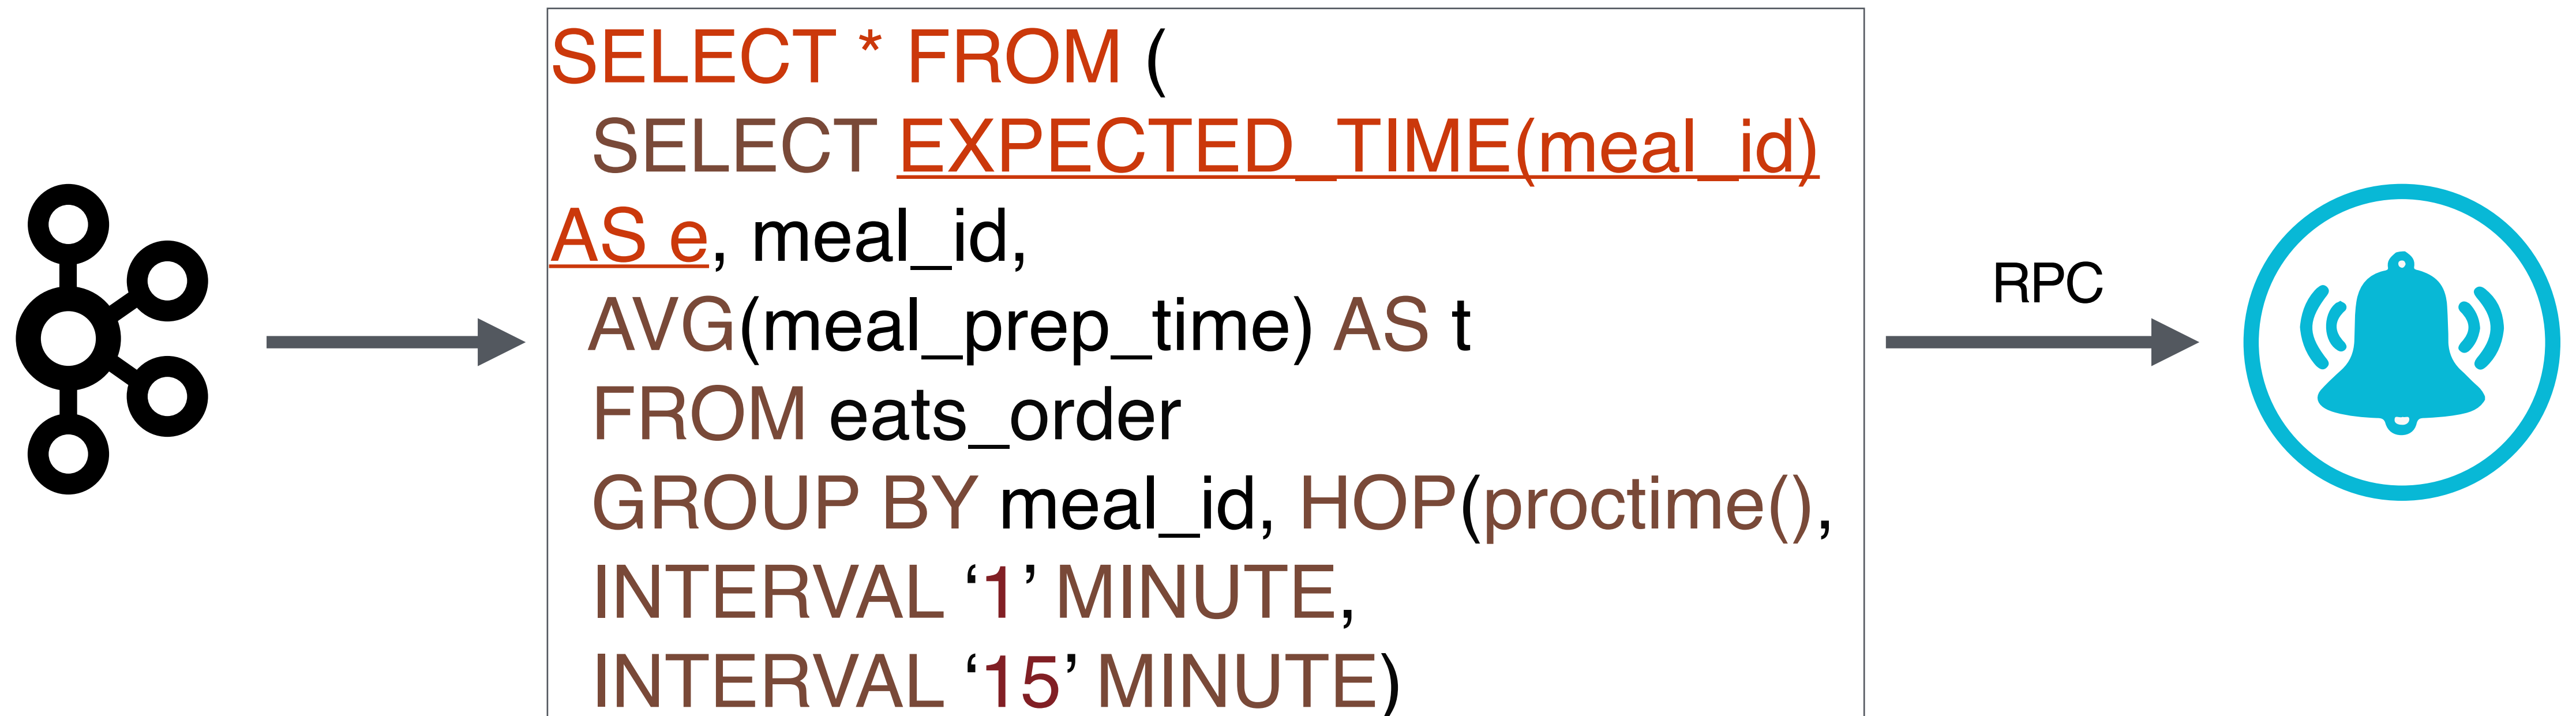
Building streaming processing applications with SQL

```
SELECT AVG(meal_prep_time) FROM  
eats_order
```

```
GROUP BY meal_id, HOP(proctime(),  
    INTERVAL '1' MINUTE,  
    INTERVAL '15' MINUTE)
```

Example (cont.)

Building streaming processing applications with SQL




Tables are more generic than analytical stores

Agenda

- Motivating example
- **Case study: ETD in UberEATS**
- Implementation
- Current status
- Conclusion


The case of UberEATS

- Three-way marketplace
- Real-time metrics
 - Estimated Time to Delivery (ETD)
- Transactions
- Demand forecasts



The case of UberEATS

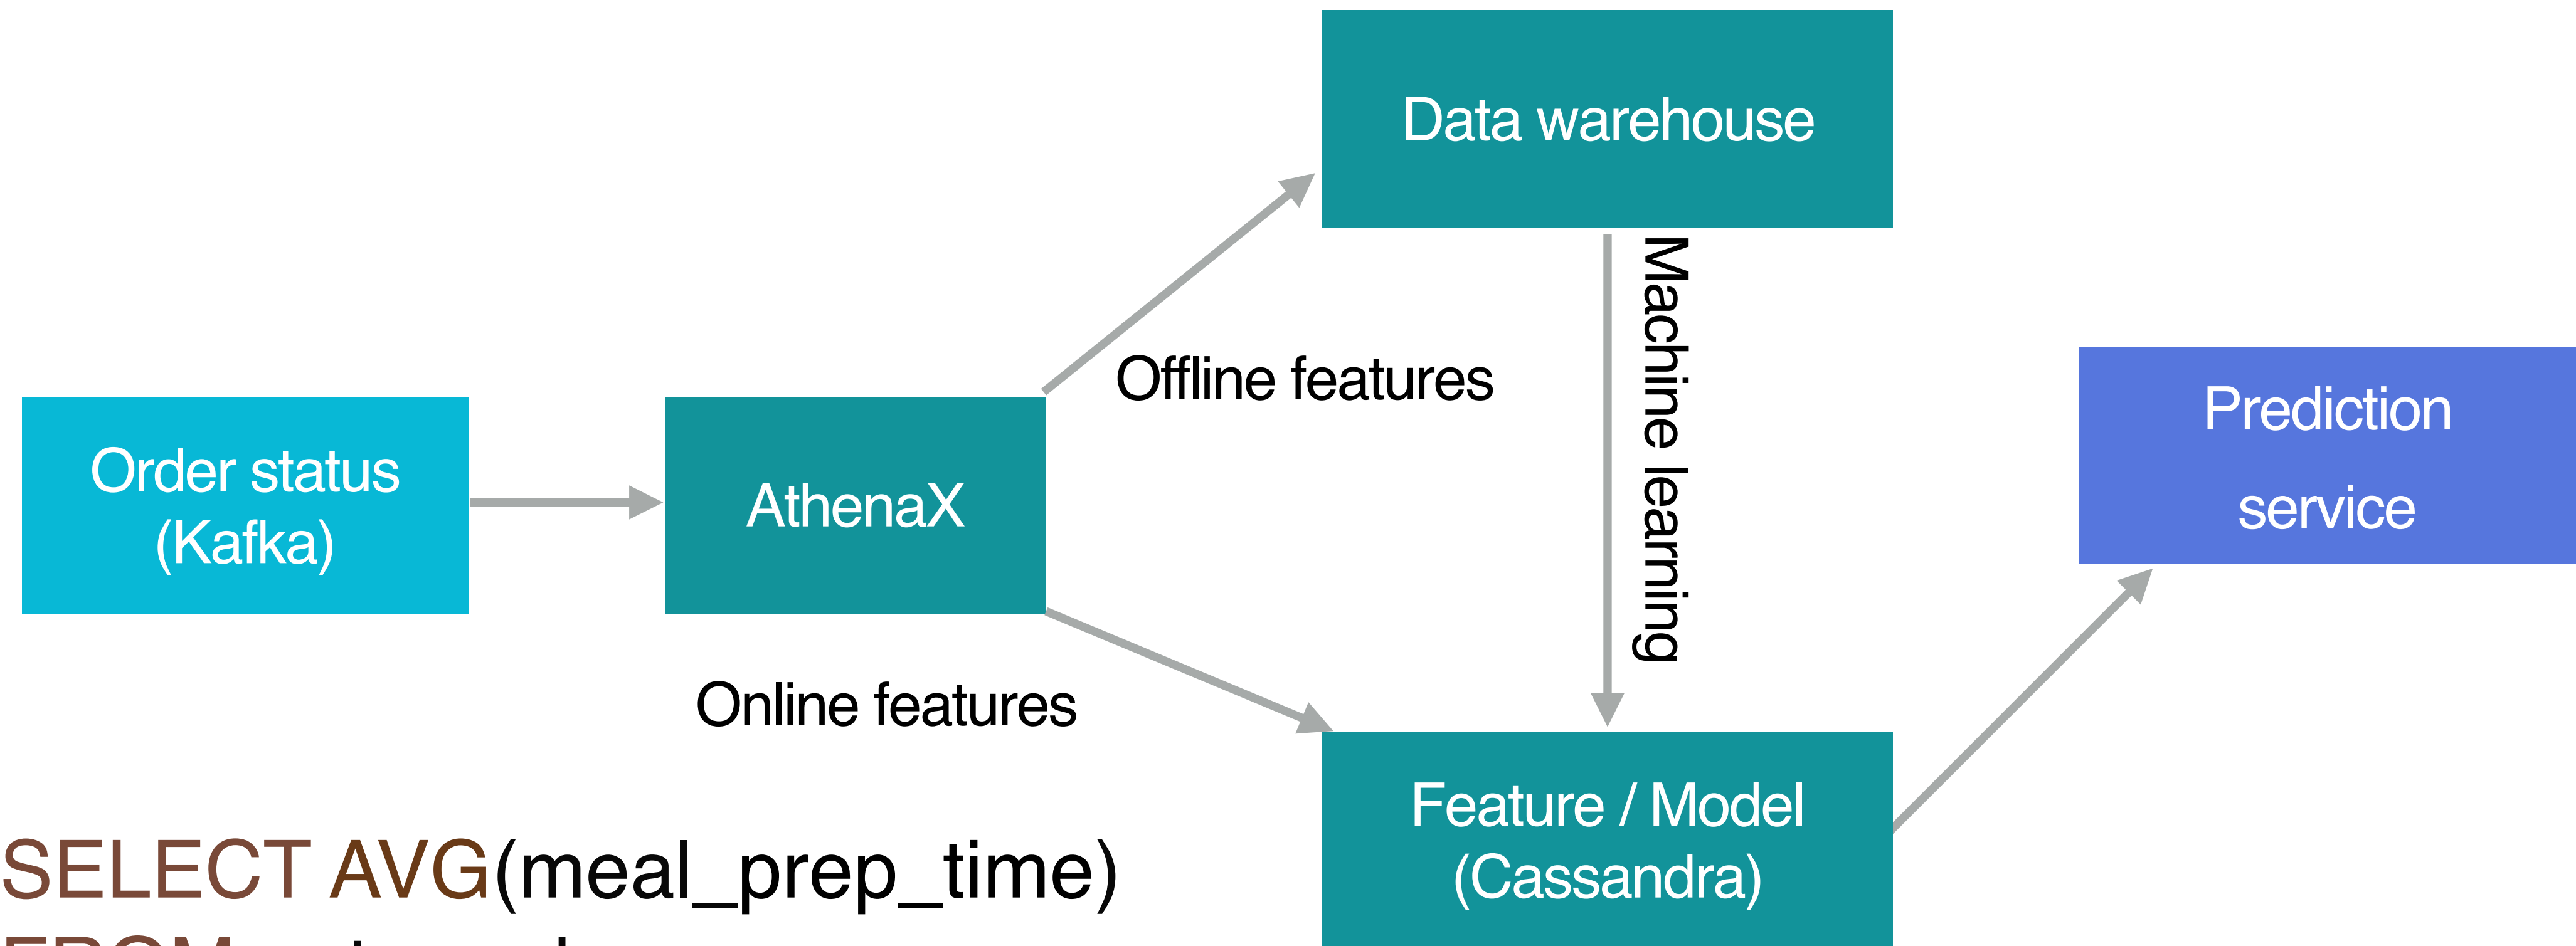
- Three-way marketplace
- Real-time metrics
 - Estimated Time to Delivery (ETD)
- Transactions
- Demand forecasts



Predicting the ETD

- Key metric: time to prepare a meal(t_{prep})
- Learn a function f : (*order status*) $\rightarrow t_{\text{prep}}$ *periodically*
- Predict the ETD for current orders using f
- AthenaX extracts features for both learnings and predictions

Architecture of the ETD service

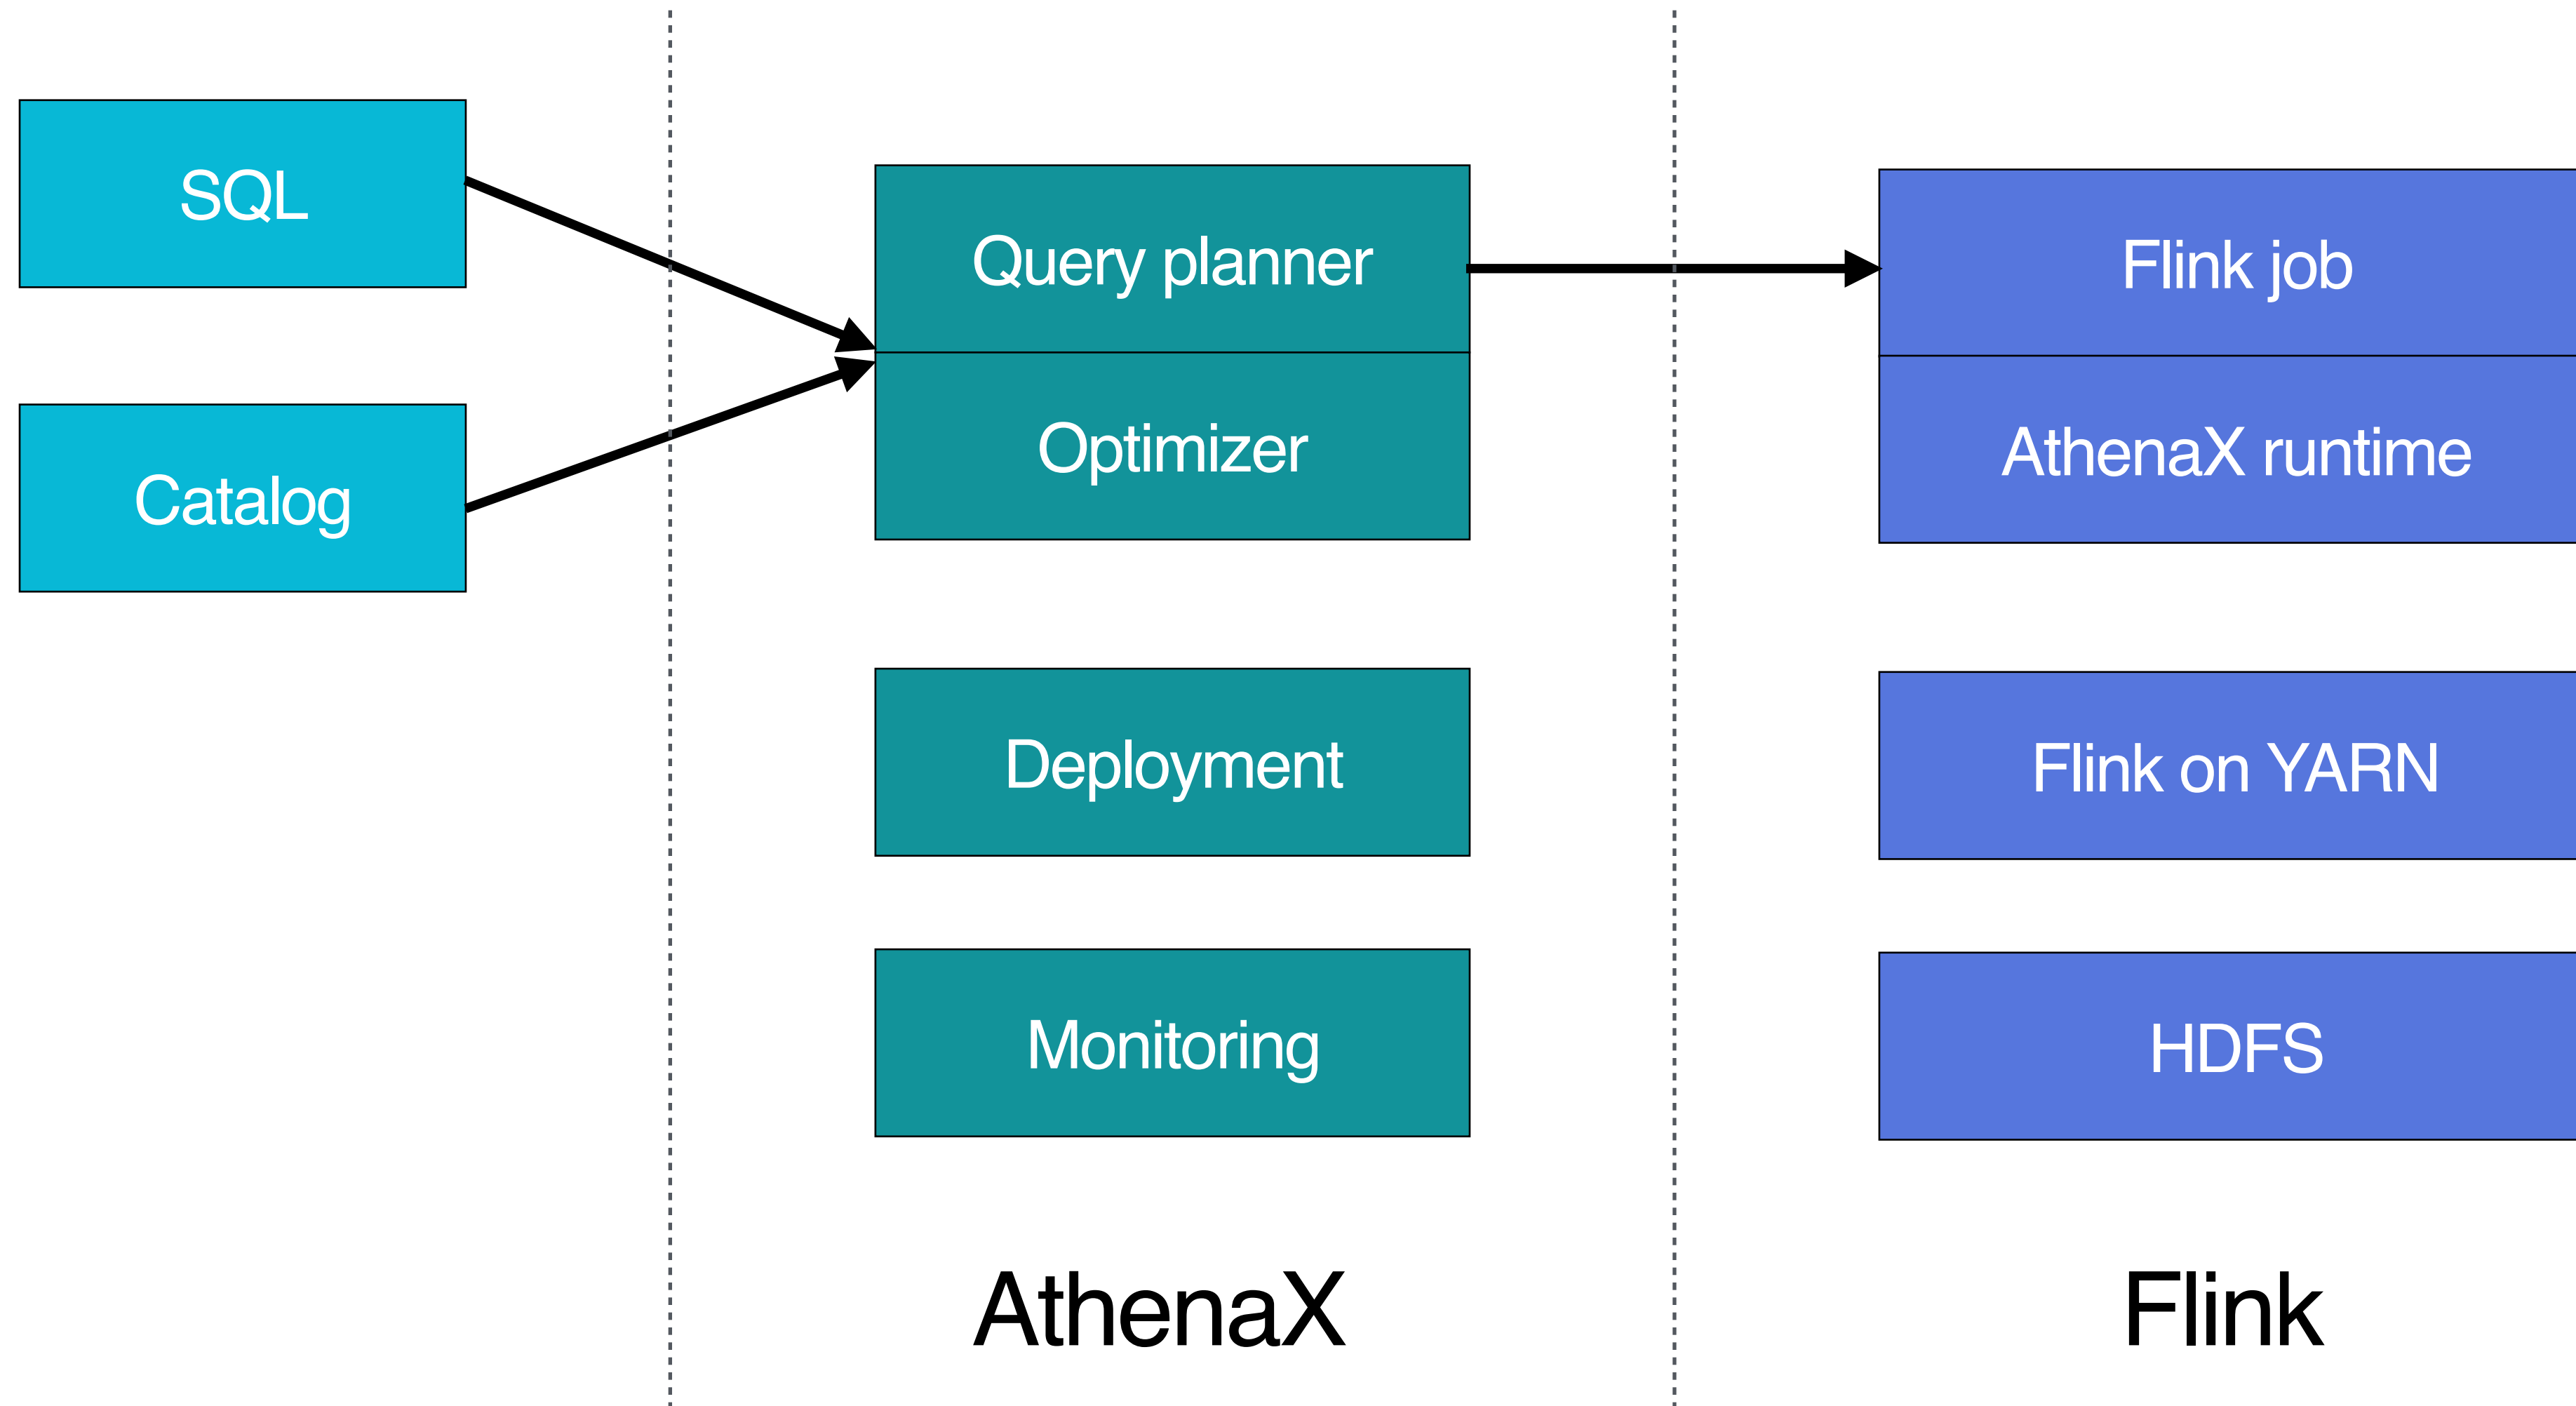


```
SELECT AVG(meal_prep_time)
FROM eats_order
GROUP BY meal_id,
HOP(proctime(),
INTERVAL '1' MINUTE,
```

Agenda

- Motivating example
- Case study: ETD in UberEATS
- **Implementation**
- Current status
- Conclusion

Architecture



Executing AthenaX applications

Compile SQLs to Flink applications

- Compilation + Code generation
 - Flink SQL APIs: SQL → Logical plans → Flink applications
 - Leverage the Volcano optimizer in Apache Calcite
- Challenges: exposing streaming semantics

Query planner

Optimizer

Deployment

Monitoring

AthenaX as a self-serving platform

Self-serving production support end-to-end

- Metadata / catalog management
- Job management
- Monitoring
- Resource management and elastic scaling
- Failure recovery

Query planner

Optimizer

Deployment

Monitoring

Agenda

- Motivating example
- Case study: ETD in UberEATS
- Implementation
- **Current status**
- Conclusion

Current status

- Pilot jobs in production
 - In the process of full-scale roll outs
- Based on Apache Flink 1.3-SNAPSHOT
 - Projection, filtering, group windows, UDF
 - Streaming joins not yet supported

Embrace the community

Contributions to the upstream

- Group window support for streaming SQL
 - CALCITE-1603, CALCITE-1615
 - FLINK-5624, FLINK-5710, FLINK-6011, FLINK-6012
- Stability fixes
 - FLINK-3679, FLINK-5631
- Table abstractions for Cassandra / JDBC (WIP)
- Available in the upcoming 1.3 release



Agenda

- Motivating example
- Case study: ETD in UberEATS
- Implementation
- Current status
- Conclusion

Conclusion

- AthenaX: write SQLs to build streaming applications
 - Treat table as a generic concept
 - Productivity: development → production in hours
- The AthenaX approach
 - SQL on streams as a platform
 - Self-serving production support end-to-end

Thank you

UBER

Compiling SQL

```
SELECT AVG(meal_prep_time)
FROM eats_order
GROUP BY meal_id,
HOP(proctime(),
INTERVAL '1' MINUTE,
```

Parsing

LogicalProject

LogicalAggregate

LogicalProject

LogicalTableScan

Planning

DataStreamCalc

DataStreamAggregate

DataStreamCalc

DataStreamScan

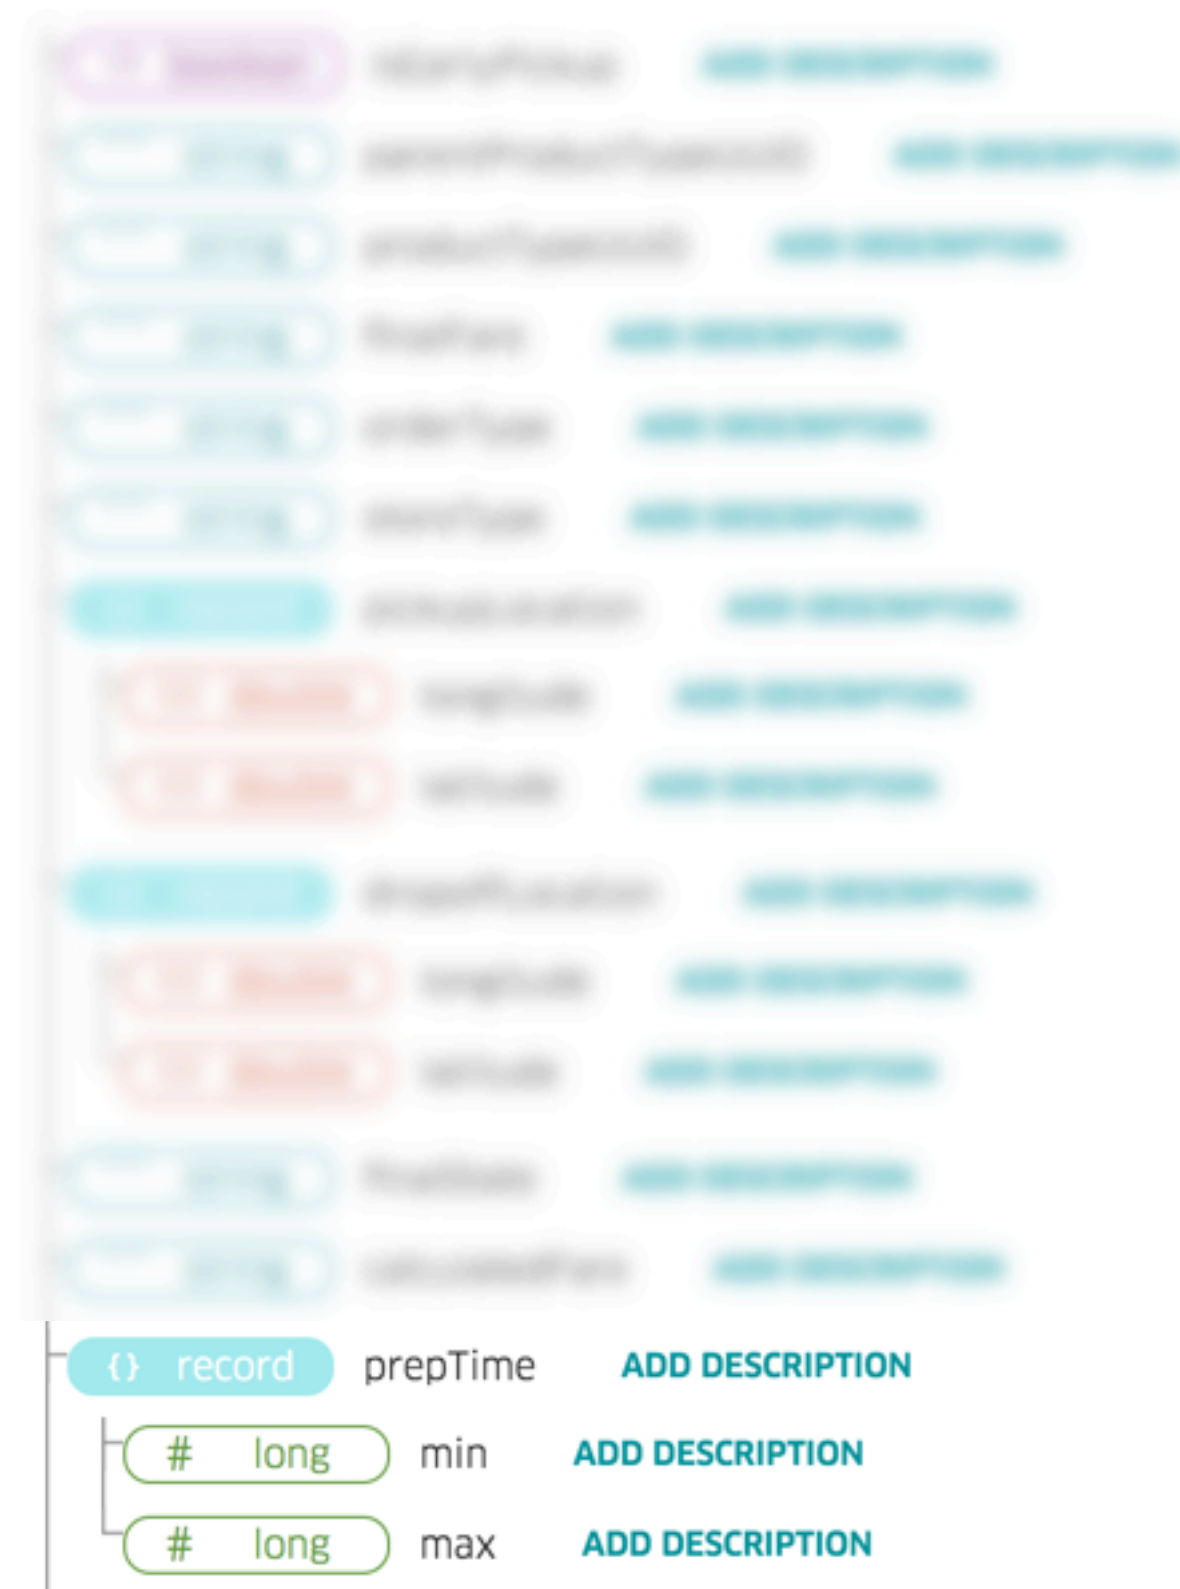


01001...

```
val eats = getEatsOrder()
eats.window(Slide
  .over("15.minutes")
  .every("1.minute"))
  .avg("meal_prep_time")
```

Lazy deserialization

Example of SQL optimization



```
SELECT  
AVG(meal_prep_time)  
FROM eats_order
```

