

Blink's Improvements to Flink SQL & Table API

FlinkForward 2017.4

Shaoxuan Wang,

Xiaowei Jiang

{xiaowei.jxw, shaoxuan.wang}
@alibaba-inc.com





About Us

■ Xiaowei Jiang

- 2014-now Alibaba
- 2010-2014 Facebook
- 2002-2010 Microsoft
- 2000-2002 Stratify

■ Shaoxuan Wang

- 2015-now Alibaba
- 2014-2015 Facebook
- 2010-2014 Broadcom



Agenda

- **Background**
- **Why SQL & Table API**
- **Blink SQL & Table API (Selected Topics)**



Background About Alibaba

■ Alibaba Group

- Operates the world's largest e-commerce platform
- Recorded GMV of \$394 Billion in year 2015, \$17.8 billion worth of GMV on Nov 11, 2016

■ Alibaba Search

- Personalized search and recommendation
- Major driver for user's traffic



[Sign In](#) | [Join Free](#)

[My Alibaba](#) ▾



AliExpress™
Smarter Shopping. Better Living!

I'm shopping for...

All Categories ▾



Products ▾

What are you looking for...

SEARCH

淘宝网
Taobao.com

宝贝 天猫 店铺



搜索



| 戳森颠覆你的生活

搜索

连衣裙夏 | 智能手机 | t恤男 | 亲子装 | 凉鞋 | 夏凉被 | 睡衣女夏 | 防晒衣 | 电风扇




What is Blink?

- **Blink – A project to make Flink work well for large scale production at Alibaba**
 - **Run on Thousands of Nodes In Production**
 - **Support Key Production Scenarios, such as Search and Recommendation**
 - **Compatible with Flink API and Ecosystem**



Documents Building for Search





Why Flink SQL & Table API

■ Unify batch and streaming

- Flink currently offers DataSet API for batch and DataStream API for streaming
- We want a single API that can run in both batch and streaming mode


■ Simplify user code

- Users only describe the semantics of their data processing
- Leave hard optimization problems to the system
- SQL is proven to be good at describing data processing
- Table API makes multi-statement data processing easy to write
- Table API also makes it possible/easy to extend standard SQL when necessary



Stream-Table Duality

Stream	
word	count
Hello	1
World	1
Hello	2
Bark	1
Hello	3



Dynamic Table	
word	count
Hello	3
World	1
Bark	1



Dynamic Tables

■ Apply Changelog Stream to Dynamic Table

- **Append Mode:** each stream record is an insert modification to the dynamic table. Hence, all records of a stream are appended to the dynamic table
- **Update Mode:** a stream record can represent an insert, update, or delete modification on the dynamic table (append mode is in fact a special case of update mode)

■ Derive Changelog Stream from Dynamic Table

- **REDO Mode:** where the stream records the new value of a modified element to redo lost changes of completed transactions
- **REDO+UNDO Mode:** where the stream records the old and the new value of a changed element to undo incomplete transactions and redo lost changes of completed transactions



Stream SQL



**Dynamic Tables generalize the concept of Static Tables
SQL serves as the unified way to describe data processing in
both batch and streaming**

There is no such thing as Stream SQL



Blink SQL & Table API

- Stream-Stream Inner Join
- User Defined Function (UDF)
- User Defined Table Function (UDTF)
- User Defined Aggregate Function (UDAGG)
- Retract (stream only)
- Over Aggregates



A Simple Query: Select and Where

id	name	price	sales	stock
1	Latte	6	1	1000
8	Mocha	8	1	800
4	Breve	5	1	200
7	Tea	4	1	2000
1	Latte	6	2	998

```
SELECT id, name, price, sales, stock  
FROM myTable WHERE name = 'Latte'
```



id	name	price	sales	stock
1	Latte	6	1	1000
1	Latte	6	2	998



Stream-Stream Inner Join

id1	name	stock
1	Latte	1000
8	Mocha	800
4	Breve	200
3	Water	5000
7	Tea	2000

id2	price	sales
1	6	1
8	8	1
9	3	1
4	5	1
7	4	1

```
SELECT id1 AS id, name, price, sales, stock  
FROM table1 INNER JOIN table2 ON id1 = id2
```



id	name	price	sales	stock
1	Latte	6	1	1000
8	Mocha	8	1	800
4	Breve	5	1	200
7	Tea	4	1	2000

*This is proposed and discussed in
FLINK-5878*



Blink SQL & Table API

- Stream-Stream Inner Join
- User Defined Function (UDF)
- User Defined Table Function (UDTF)
- User Defined Aggregate Function (UDAGG)
- Retract (stream only)
- Over Aggregates



User Defined Function (UDF)

Create and use a UDF is very simple and easy:

```
object AddFunc extends ScalarFunction {  
  def eval(a: Long, b: Long): Long = a + b  
  @varargs  
  def eval(a: Int*): Int = a.sum  
}  
  
tEnv.registerTable("MyTable", table)  
tEnv.registerFunction("addFunc", AddFunc)  
val sqlQuery =  
  "SELECT addFunc(long1,long2), addFunc(int1,int2,int3) FROM MyTable"
```

We recently have enhanced UDF/UDTF to let them support variable types and variable arguments (FLINK-5826)



User Defined Table Function (UDTF)

line
Tom#23 Jark#17 David#50

Scalar → Table (multi rows and columns)



name	age
Tom	23
Jack	17
David	50

```
SELECT name, age
FROM myTable,
  LATERAL TABLE(splitFunc(line))
AS T(name, age)
```

```
case class User(name: String, age: Int)
class SplitFunc extends TableFunction[User] {
  def eval(str: String): Unit = {
    str.split(" ").foreach{ e =>
      val subSplits = e.split("#")
      collect(User(subSplits(0), subSplits(1).toInt))
    }
  }
}
```

We have shipped UDTF in flink 1.2 (FLINK-4469).



User Defined Aggregate Function (UDAGG) - Motivation

Flink has built-in aggregates (count, sum, avg, min, max) for SQL and table API

```
SELECT name, SUM(price), COUNT(sales),  
        MAX(price), MIN(price), AVG(price)  
FROM myTable  
GROUP BY name
```

What if user wants an aggregate that is not covered by built-in aggregates, say a weighted average aggregate?

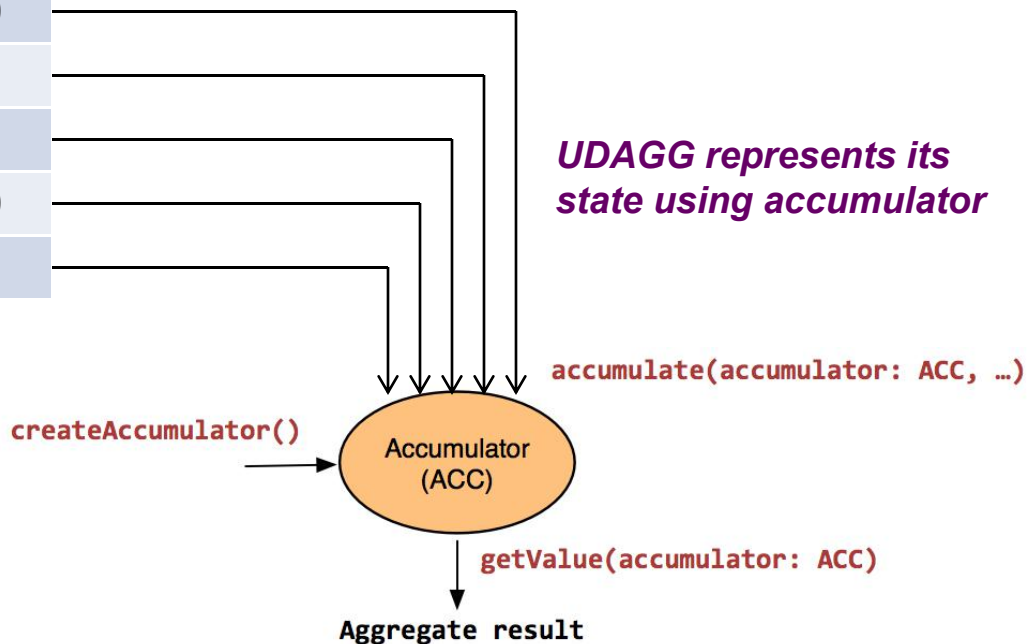
We need an aggregate interface to support user defined aggregate function.



UDAGG – Accumulator (ACC)

id	name	price	sales	stock
1	Latte	6	1	1000
8	Mocha	8	1	800
4	Breve	5	1	200
7	Tea	4	1	2000
1	Latte	6	2	998

```
SELECT name, SUM(price), COUNT(sales),  
       MAX(price), MIN(price), AVG(price)  
FROM myTable  
GROUP BY name
```





UDAGG – Interface

UDAGG Interface

```
Abstract class AggregateFunction[T, ACC] extends UserDefinedFunction {  
  def createAccumulator(): ACC  
  def getValue(accumulator: ACC): T  
}  
  
/* The implementations of accumulate must be declared publicly,  
not static and named exactly as "accumulate". accumulate method  
can be overloaded */  
def Accumulate(accumulator: ACC, [user defined inputs]): Unit
```

SQL Query

```
SELECT type, weightAvgFun(weight, cnt)  
FROM myTable  
GROUP BY type
```

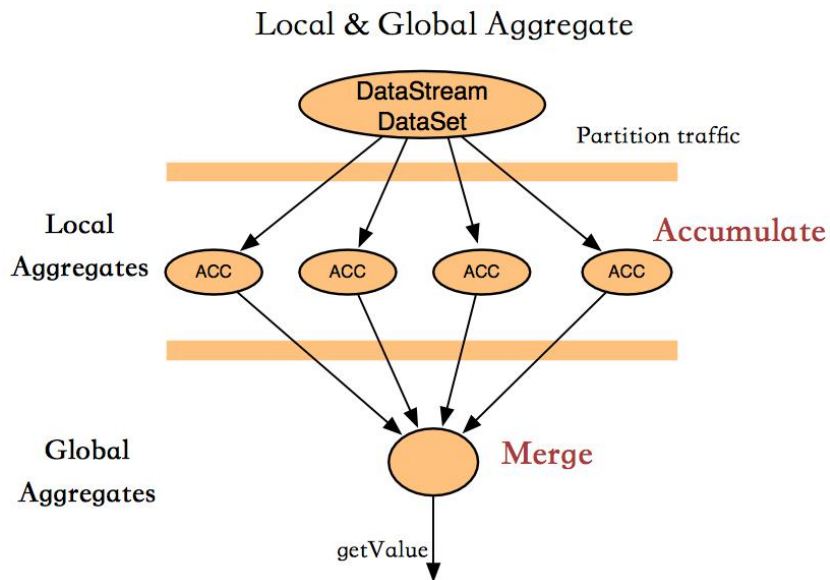
UDAGG example: a weighted average

```
public static class WeightedAvgAccum {  
  public long sum = 0; public int count = 0; }  
public static class WeightedAvg  
  extends AggregateFunction<Long, WeightedAvgAccum> {  
  @Override  
  public WeightedAvgAccum createAccumulator() {  
    return new WeightedAvgAccum(); }  
  @Override  
  public Long getValue(WeightedAvgAccum accumulator) {  
    if (accumulator.count == 0) return null;  
    else return accumulator.sum/accumulator.count; }  
  public void accumulate(  
    WeightedAvgAccum accumulator, long iValue, int iWeight) {  
    accumulator.sum += iValue * iWeight;  
    accumulator.count += iWeight; }  
}
```



UDAGG – Merge

How to count the total visits on TaoBao web pages in real time?



```
Abstract class AggregateFunction[T, ACC] extends UserDefinedFunction {  
  def createAccumulator(): ACC  
  def getValue(accumulator: ACC): T  
}  
/* The implementations of following methods: accumulate(MUST  
have), merge(OPTIONAL) must be declared publicly, not static and  
named exactly as "accumulate", "merge"*/  
  def Accumulate(accumulator: ACC, [user defined inputs]): Unit  
  def merge(it: java.util.iterator[ACC]): ACC
```

Motivated by local & global aggregate (and session window merge etc.), we need a merge method which can merge the partial aggregated accumulator into one single accumulator



UDAGG - Retract - Motivations

```
SELECT
cnt,
COUNT(word) AS freq
FROM (
  SELECT word, COUNT(num) AS cnt
  FROM Table GROUP BY word
) GROUP BY cnt
```

Stream

word	num
Hello	1
World	1
Bark	1

Keyed Table

word	cnt
Hello	1
World	1
Bark	1

Keyed Table

cnt	freq
1	3

Stream

word	num
Hello	1
World	1
Bark	1
Hello	1

source table
adds a new Row
(hello,1)

Keyed Table

word	cnt
Hello	1 -> 2
World	1
Bark	1

Without Retraction

(hello,2)

Keyed Table

cnt	freq
1	3
2	1

**Incorrect! This value
should be 2**



UDAGG - Retract – Motivations

Stream

word	num
Hello	1
World	1
Bark	1
Hello	1

source table
adds a new Row

(hello,1)



Keyed Table

word	cnt
Hello	1 -> 2
World	1
Bark	1

Without Retraction

(Hello,2,accumulate)



Incorrect! this value
should be 2

Keyed Table

cnt	freq
1	3
2	1

With Retraction

(Hello,1,retract)

(Hello,2,accumulate)



Keyed Table

cnt	freq
1	2
2	1

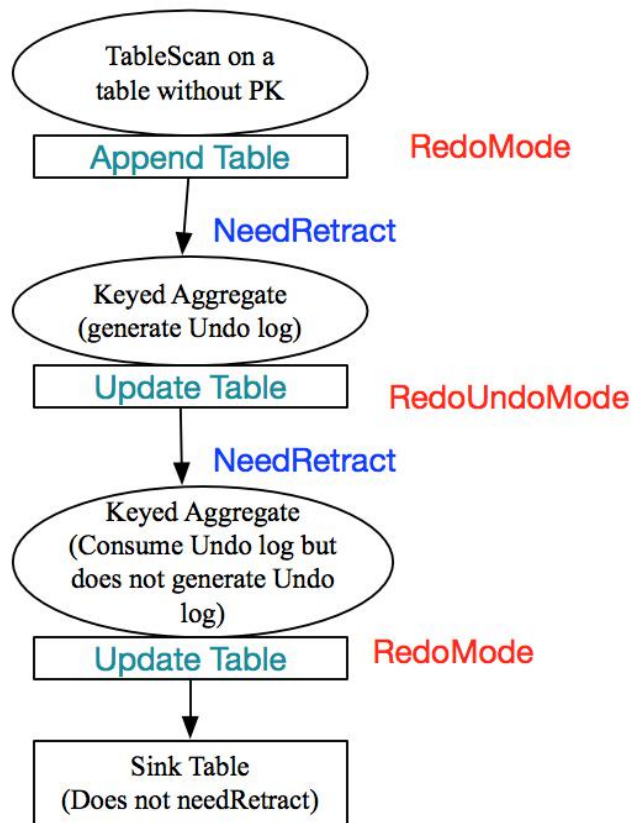
We need a retract method in UDAGG, which can retract the input values from the accumulator



Retract – Solution

- Retraction is introduced to handle updates
- We use query optimizer to decide where the retraction is needed.

The design doc and the progress of retract implementation are tracked in [FLINK-6047](#). A FLIP for retract is on the way. We aim to release it in flink-1.3





UDAGG – Summary

```
Abstract class AggregateFunction[T, ACC] extends UserDefinedFunction {  
  def createAccumulator(): ACC  
  def getValue(accumulator: ACC): T  
}  
  
/* The implementations of following methods: accumulate(MUST have),  
merge(OPTIONAL), and retract(OPTIONAL) must be declared publicly, not  
static and named exactly as "accumulate", "merge", and "retract" */  
def accumulate(accumulator: ACC, [user defined inputs]): Unit  
def merge(it: java.util.iterator[ACC]): ACC  
def retract(accumulator: ACC, [user defined inputs]): Unit
```

Master JIRA for UDAGG is [FLINK-5564](#). We plan to ship it in release 1.3.



Blink SQL & Table API

- Stream-Stream Inner Join
- User Defined Function (UDF)
- User Defined Table Function (UDTF)
- User Defined Aggregate Function (UDAGG)
- Retract (stream only)
- Over Aggregates



Over Aggregates

Calculate moving average (in the past 5 seconds), and emit the result for each record

time	itemID	price
1000	101	1
3000	201	2
4000	301	3
5000	101	1
5000	401	4
7000	301	3
8000	501	5
10000	101	1

```
SELECT
  time, itemID, MovingAverage(price) OVER
  (
    ORDER BY RowTime()
    RANGE
      BETWEEN INTERVAL '5' SECOND PRECEDING
      AND CURRENT ROW
  ) AS avgPrice
FROM myTable
```

time	itemID	avgPrice
1000	101	1
3000	201	1.5
4000	301	2
5000	101	2.2
5000	401	2.2
7000	301	2.6
8000	401	3
10000	101	2.8

time based Group Aggregate is not able to differentiate two records with the same row time.



Group/Over Aggregates

- **Grouping methods:** Groupby / Over
- **Window types:**
 - Time/Count + TUMBLE/SESSION/SLIDE window;
 - OVER Range/Rows window
- **Time types:** Event time; Process time (only for stream)

We have been working closely with team dataArtisans, from the design to the implementation on FLIP11 (FLINK-4557). Upon now, except the unbounded group aggregate, all other group/over aggregates are fully supported via SQL query. We are working on the support for table API.



Current Status of Flink SQL & Table API

- Flink blog: “Continuous Queries on Dynamic Tables” (*posted at <https://flink.apache.org/news/2017/04/04/dynamic-tables.html>*)
- UDF (*several improvements will be released in 1.3*)
- UDTF (*FLINK-4469, released in 1.2*)
- UDAGG (*FLINK-5564, target for release 1.3*)
- Group/Over Window Aggregate (*FLINK-4557, target for release 1.3*)
- Retract (*FLINK-6047, target for release 1.3*)
- Unbounded Stream Group Aggregate (*FLINK-6216, bundled with retract design*)
- Stream-Stream Inner Join (*FLINK-5878, TBD*)

We will keep merging Blink SQL & Table API back to Flink



Q & A

Thank You!