

Dynamically Configured Stream Processing Using Flink & Kafka

David Hardwick
Sean Hester
David Brelloch

<https://github.com/brelloch/FlinkForward2017>

Multi-SaaS Management



Confluence



Namely

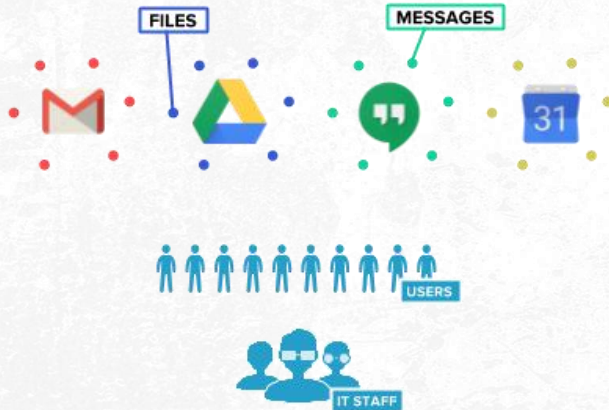


“What does that mean?”

Complex Interconnect of Apps, People and Data

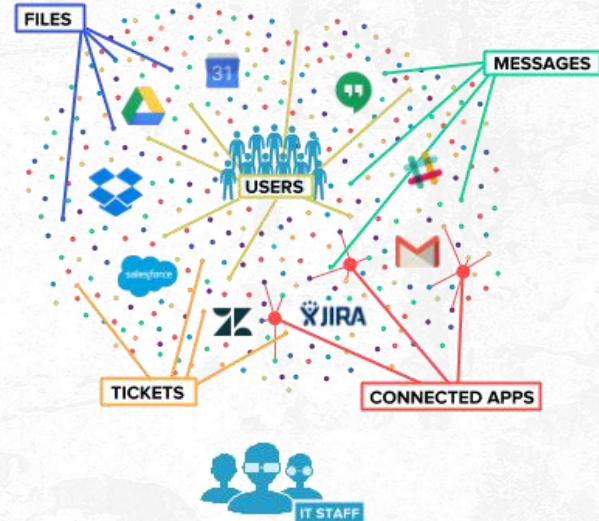
A Few Years Ago

Cloud Office

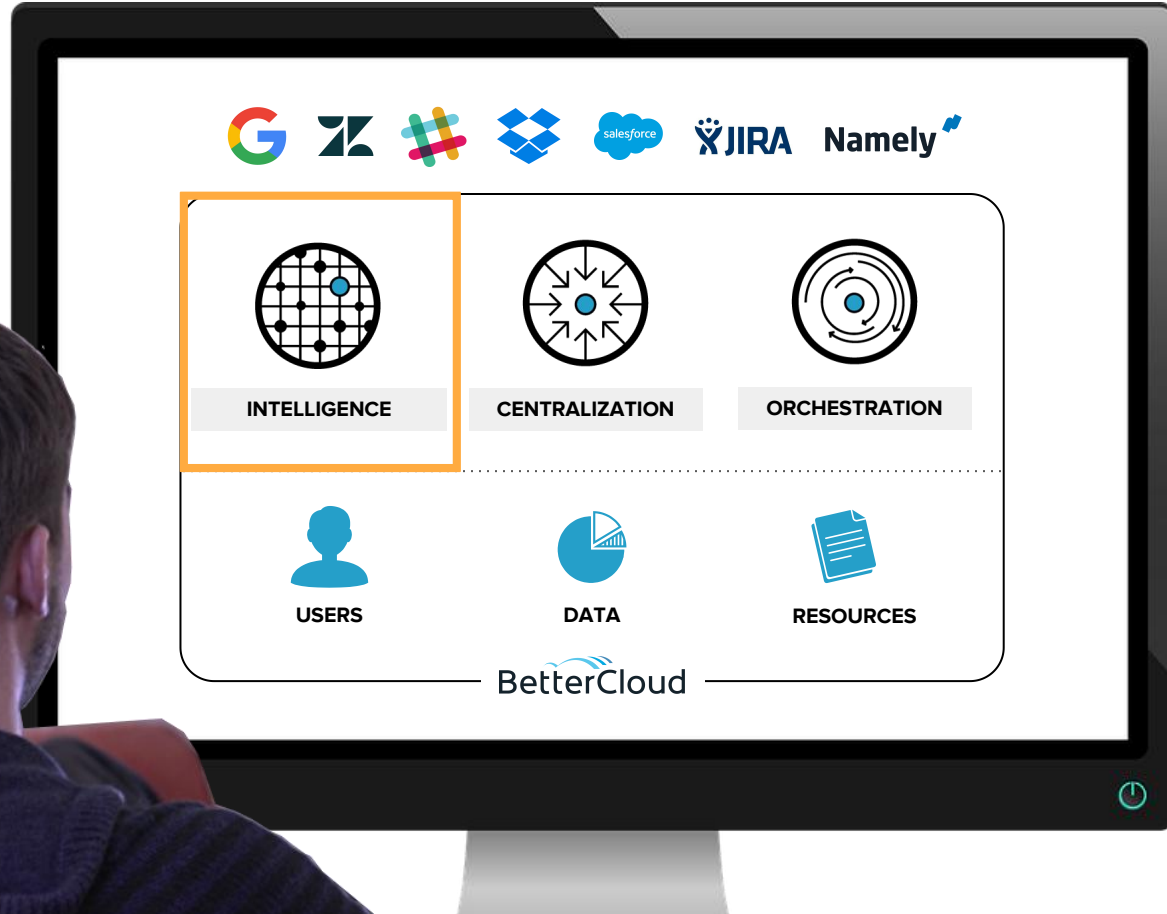


This is Now

SaaS Sprawl



“Multi-SaaS Management” is like **Mission Control** for IT Admins



We process 100s millions of messages daily

CONNECTORS

CENTRALIZE

INTELLIGENCE

ORCHESTRATION

PUBLIC APIs



Data
Ingestion

Normalization

Role-Based Privileges

USERS

Internal Users

External Users

Memberships

DATA

Files

Connected
Apps

RESOURCES

Calendars

Devices

Conference Rooms

Alert
Detections

Single
or Bulk
Actions

Policy
Violations

Automated
Workflows




Multi-SaaS Management

BetterCloud Platform

CONNECTORS



CENTRALIZE



Role-Based Privileges



USERS

Internal Users
External Users
Memberships



DATA

Files
Connected Apps




RESOURCES

Calendars
Devices
Conference Rooms

INTELLIGENCE



ORCHESTRATION



BetterCloud Platform

Multi-SaaS Management

BetterCloud Platform

CONNECTORS

CENTRALIZE

INTELLIGENCE

ORCHESTRATION

PUBLIC APIs



Data
Ingestion

Normalization

Role-Based Privileges

USERS

Internal Users
External Users
Memberships

DATA

Files
Connected
Apps

RESOURCES

Calendars
Devices
Conference Rooms

Alert
Detections

Single
or Bulk
Actions

Policy
Violations

Automated
Workflows



BetterCloud Platform

NAME ↑	EMAIL	APPLICATION	STATUS	ROLE
Albert (AJ) LeGaye	albert.legaye@demobettercloud.com	G	Active	Admin
Alex Miller	alex.miller@demobettercloud.com	G #	Active	Admin
Amanda McCarthy	amanda.mccarthy@demobettercloud.com	G #	Active	Admin
Andrew McGonnigle	andrew.mcgonnigle@demobettercloud.com	G	Active	Admin
Ben Howard	ben.howard@demobettercloud.com	G #	Active	Admin
BetterCloud Notifications	notifications@demobettercloud.com	G	Active	End User
Bill Quinn	bill.quinn@demobettercloud.com	G	Active	Admin
Caitlin McDevitt	caitlin.mcdevitt@demobettercloud.com	G	Active	Admin
Caroline Thompson	caroline.thompson@demobettercloud.com	G #	Active	Admin
Chris Test	chris.test@demobettercloud.com	G	Active	End User
Christopher		G #		

ALERTS

DIRECTORY

USERS

GROUPS

DATA MANAGEMENT

WORKFLOWS

CONNECTORS

PRIVILEGES

AUDIT LOGS

Privacy Policy

NAME

Albert

Alex

Ama

McCa

Andr

McGo

Ben

Bette

Notif

Bill Q

Caitl

Caro

Thon

Chris

Chris

Chris

Chris

Chris

Chris

**Alex Miller**

Sales Representative



BetterCloud



Slack



Zendesk



Google



Dropbox

Summary

Memberships

Files

GROUP NAME ↑	APPLICATION	ROLE	TYPE
admin_alerts		Member	group
ATL		—	group
bctest		—	group
BetterCloud Soccer Team		—	group
Customer List		—	group
Everyone at Demo BetterCloud		Member	group
general		Member	group
Human Resources		—	group

1 of 2

View 10 per page.

Viewing 1 - 10 of 14

352

FILES

5

CALENDARS

4

INTEGRATIONS

32

GROUPS

12%

STORAGE USED

Live Chat

data centralization

Multi-SaaS Management

BetterCloud Platform

CONNECTORS

CENTRALIZE

INTELLIGENCE

ORCHESTRATION

PUBLIC APIs



Data
Ingestion

Normalization

Role-Based Privileges

USERS

Internal Users
External Users
Memberships

DATA

Files
Connected
Apps

RESOURCES

Calendars
Devices
Conference Rooms

Alert
Detections

Single
or Bulk
Actions

Policy
Violations

Automated
Workflows

BetterCloud Platform

Multi-SaaS Management

BetterCloud Platform

CONNECTORS

PUBLIC APIs



Data
Ingestion

Normalization

Role-Based Privileges

USERS

Internal Users
External Users
Memberships

DATA

Files
Connected
Apps

RESOURCES

Calendars
Devices
Conference Rooms

INTELLIGENCE

Alert
Detections

Policy
Violations

ORCHESTRATION

Single
or Bulk
Actions

Automated
Workflows

BetterCloud Platform

ALERTS

DIRECTORY

DATA MANAGEMENT

WORKFLOWS

CONNECTORS

PRIVILEGES

AUDIT LOGS

Users with No Depart

Empty Groups

Super Administrator

Super Administrator

Super Administrator

Users without Two-F

Administrator Count

Users with No Name

Users with No Title

Administrator Adde

Users with No Phone

Users without Two-Factor Authentication

Details

Severity: Major

Application:  Slack

Date Occurred: 06/Feb/2017 10:35 PM

Description: Users who do not have two-factor authentication set up

Threshold: 1

Count: 10

USER EMAIL

DATE/TIME TRIGGERED

michael.matmon@demobettercloud.com

06/Feb/2017 10:35 PM

caroline.thompson@demobettercloud.com

17/Jan/2017 5:40 PM

taylor.gould@demobettercloud.com

17/Dec/2016 6:29 PM



Live Chat

operational intelligence

Multi-SaaS Management

BetterCloud Platform

CONNECTORS

PUBLIC APIs



Data
Ingestion

CENTRALIZE

Normalization

Role-Based Privileges



USERS

Internal Users
External Users
Memberships



DATA

Files
Connected
Apps



RESOURCES

Calendars
Devices
Conference Rooms

INTELLIGENCE

Alert
Detections

Policy
Violations

ORCHESTRATION

Single
or Bulk
Actions

Automated
Workflows



BetterCloud Platform

Multi-SaaS Management

BetterCloud Platform

CONNECTORS

PUBLIC APIs



Data
Ingestion

Normalization

Role-Based Privileges

USERS

Internal Users
External Users
Memberships

DATA

Files
Connected
Apps

RESOURCES

Calendars
Devices
Conference Rooms

INTELLIGENCE

Alert
Detections

Policy
Violations

ORCHESTRATION

Single
or Bulk
Actions

Automated
Workflows

BetterCloud Platform

- ALERTS
- DIRECTORY
- USERS
- GROUPS
- DATA MANAGEMENT
- WORKFLOWS
- CONNECTORS
- PRIVILEGES
- AUDIT LOGS

	NAME ↑	EMAIL	APPLI
<input type="checkbox"/>	asdfasdf3 asdfasdf	aa_asdfasdf3.asdfasdf-bb@0000ff.co	G
<input checked="" type="checkbox"/>	brenda harden	brenda@0000ff.co	G
<input checked="" type="checkbox"/>	Brenda Harden	brenda.harden@0000ff.co	Z
<input type="checkbox"/>	Brian Test	briantest@0000ff.co	G
<input checked="" type="checkbox"/>	Christina Kang	christina.kang@0000ff.co	G
<input checked="" type="checkbox"/>	David Hardwick	david@0000ff.co	G
<input type="checkbox"/>	directory user	directory.user@0000ff.co	G
<input checked="" type="checkbox"/>	Jason Lang	jason@0000ff.co	G
<input type="checkbox"/>	john.mcintosh	john.mcintosh@0000ff.co	#
<input type="checkbox"/>	Johnny McIntosh	mcintosh.john@0000ff.co	G Z
<input type="checkbox"/>	kevin willow	kevin.willow@0000ff.co	G

All

Slack

Zendesk

Dropbox

Google

Slack

Send Integration Logs

Set User's Phone Number

Send Direct Message

Remove from Channel

Add To Channel

Create Reminder

Remove From User Group

Add To User Group

Set User's Title

Zendesk

Sign Out User

Reset Password

Suspend User

Delete User

Hint: type a to quickly open and search this menu

close

Flink detects and evaluates

ACTIVE WORKFLOW

This workflow is active. Changes to this workflow will be reflected the next time it runs.

CANCEL

SAVE

LIBRARY

WHEN

IF

THEN

ALL			
		Add To Channel	+
Dropbox	>	Add To Group	+
Google	>	Add To User Group	+
Slack	>	Archive Group	+
Zendesk	>	Assign License	+
		Block Mobile Device	+
		Change Role in Group	+
		Copy Group Membership	+

Add To Channel

This action adds a user to a public Slack channel.

1. Deprovisioning Users

23 / 50

THEN

1. **HIDE** user [target user](#) in directory
2. Delete app-specific passwords for user [target user](#)
3. Remove user [target user](#) from [all groups](#) group(s)
4. Delete 2-step verification backup codes for user [target user](#)
5. Transfer ownership of all files for [target user](#) to [sharedinbox7@demobettercloud.com](#)
6. Suspend user [target user](#)
7. Move user [target user](#) to [Deprovision - Complete](#) Org Unit

PROPERTIES

Live Chat

action orchestration

Heavy use of Flink here



CONNECTORS

CENTRALIZE

OBJECTS

INTELLIGENCE

ORCHESTRATION

PUBLIC APIs



Data
Ingestion

Normalization

Role-Based Privileges



USERS

Internal Users

External Users

Memberships



DATA

Files

Connected
Apps



RESOURCES

Calendars

Devices

Conference Rooms

Alert
Detections

Policy
Violations

Single
or Bulk
Actions

Automated
Workflows



BetterCloud Platform

Business Needs and Challenges

- Alert quickly < 2s latency from ingest to alert
- Scale tremendously ~100M events/day ...and growing
- Keep innovating daily deployments to production
- Show the power detect global alerts across all customers
- Let customers drive custom alerts for every customer
- Replay history “go back in time” for new alerts ...or alert bugs

“Okay, can you go deeper?”

Nope, but Sean and David can!

Foundational Prior Learnings

1. Event Stream Processing is a better model for Operational Intelligence than Ingest + Query

Foundational Prior Learnings

1. Event Stream Processing is a better model for Operational Intelligence than Ingest + Query
2. Rules Engines add significantly more value than hard-coded rules

We process 100s of millions of messages daily

CONNECTORS

CENTRALIZE

INTELLIGENCE

ORCHESTRATION

PUBLIC APIs



Data
Ingestion

Normalization

Role-Based Privileges

USERS

Internal Users

External Users

Memberships

DATA

Files

Connected
Apps

RESOURCES

Calendars

Devices

Conference Rooms

Alert
Detections

Single
or Bulk
Actions

Policy
Violations

Automated
Workflows



BetterCloud Platform

Foundational Prior Learnings

1. Event Stream Processing is a better model for Operational Intelligence than Ingest + Query
2. Rules Engines add significantly more value than hard-coded rules

- Jayway JsonPath: <https://github.com/jayway/JsonPath>
- Executes queries against JSON documents
- For us, simplifies the creation of a rules interface for our non-technical users

JsonPath Continued

```
1 {  
2   "store": {  
3     "book": [  
4       {  
5         "category": "reference",  
6         "author": "Nigel Rees",  
7         "title": "Sayings of the Century",  
8         "price": 8.95  
9       },  
10      {  
11        "category": "fiction",  
12        "author": "Evelyn Waugh",  
13        "title": "Sword of Honour",  
14        "price": 12.99  
15      }  
16    ],  
17    "bicycle": {  
18      "color": "red",  
19      "price": 19.95  
20    }  
21  },  
22  "expensive": 10  
23 }
```

Example Paths and Output:

`$.expensive => 10`

`$.store.bicycle.price => 19.95`

`$.store.book[0].author => "Nigel Rees"`

`$.store.book.length > 0 => true`

`$.store.book[?(@.category ==
 'fiction')].price > 10.00 => { ... }`

JsonPath Wrapped in a User Interface

ALERT x EDIT ALERT

Drop

Goog

Slack

Alert Details

Title *
Team Admin Added 16 / 100

Description *
Users that have had admin permissions added 43 / 500

Alert Trigger Conditions

When a...

Application *
Dropbox

Object *
User

Event Type *
Team Admin Added Events

AND

Data Point *
Is Active

Operator *
equals

Operand *
true

X

AND

Data Point *
Manager

Operator *
is empty

X

ADD CONDITION

Timing & Thresholds

Occurs...

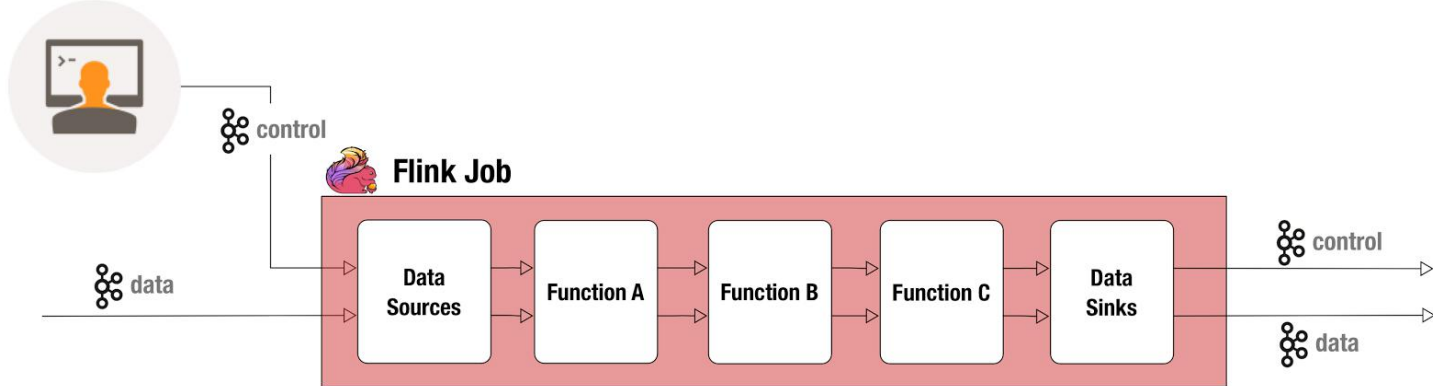
Time Period
When total count Reached

Operator
Is Greater Than

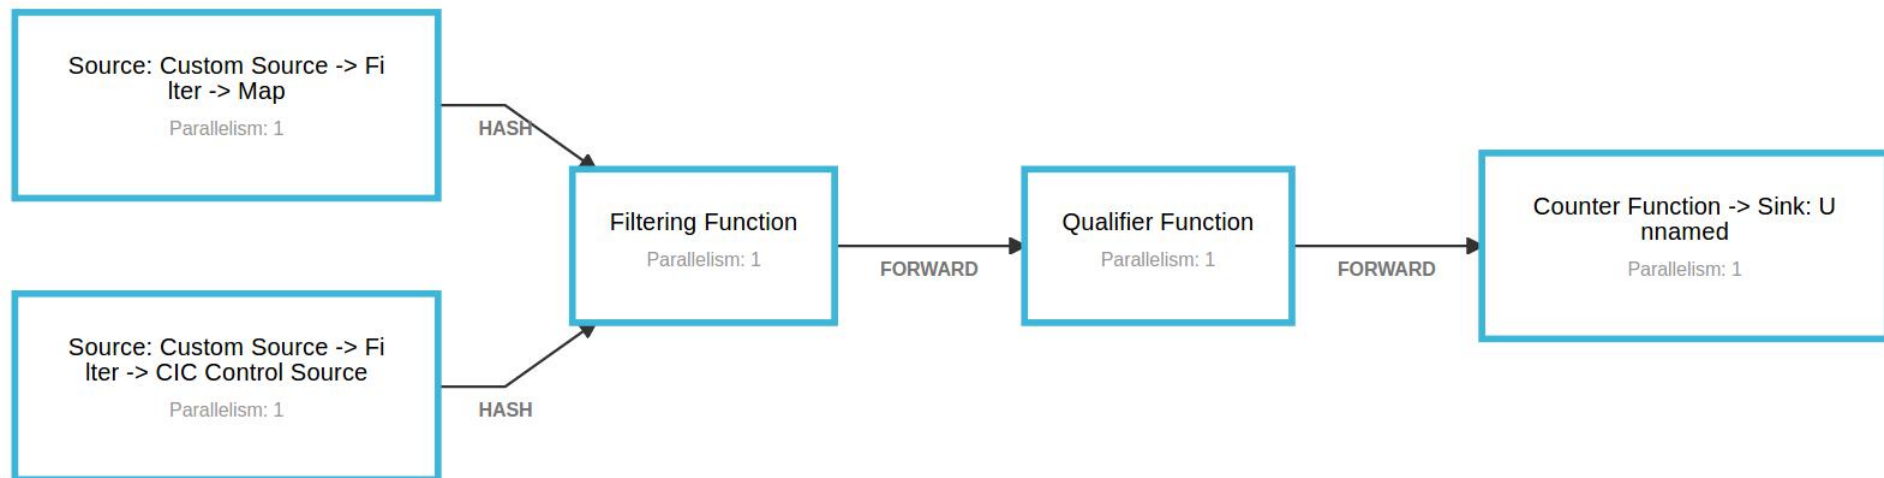
Threshold *
1

The Control Stream

User or Program Generated



The Solution



Models

```
case class CustomerEvent(customerId: UUID, payload: String)
```

```
case class ControlEvent(customerId: UUID, alertId: UUID, alertName: String, alertDescription:  
String, threshold: Int, jsonPath: String, bootstrapCustomerId: UUID)
```

Sources - Event Stream

```
val eventStream = env.addSource(new FlinkKafkaConsumer09("events", new CustomerEventSchema(),  
properties))  
    .filter(x => x.isDefined)  
    .map(x => x.get)  
    .keyBy((ce: CustomerEvent) => { ce.customerId } )
```

Sources - Control Stream

```
val controlStream = env.addSource(new FlinkKafkaConsumer09("controls", new ControlEventSchema(),
properties))

    .filter(x => x.isDefined)
    .map(x => x.get)
    .name("Control Source")
    .split((ce: ControlEvent) => {
        ce.customerId match {
            case Constants.GLOBAL_CUSTOMER_ID => List("global")
            case _ => List("specific")
        }
    })
```

Sources - Control Stream Continued

```
val globalControlStream = controlStream.select("global").broadcast
```

```
val specificControlStream = controlStream.select("specific")  
    .keyBy((ce: ControlEvent) => { ce.customerId })
```


Sources - Join the Streams

```
// Join the control and event streams
```

```
val filterStream = globalControlStream.union(specificControlStream)  
    .connect(  
        eventStream  
    )
```

Work on Streams - Filtering

```
class FilterFunction() extends RichCoFlatMapFunction[ControlEvent, CustomerEvent, FilteredEvent] {  
  var configs = new mutable.ListBuffer[ControlEvent]()  
  
  override def flatMap1(value: ControlEvent, out: Collector[FilteredEvent]): Unit = {  
    configs = configs.filter(x => (x.customerId != value.customerId) && (x.alertId != value.alertId))  
    configs.append(value)  
  }  
  
  override def flatMap2(value: CustomerEvent, out: Collector[FilteredEvent]): Unit = {  
    val eventConfigs = configs.filter(x => (x.customerId == value.customerId) || (x.customerId ==  
Constants.GLOBAL_CUSTOMER_ID))  
    if (eventConfigs.size > 0) {  
      out.collect(FilteredEvent(value, eventConfigs.toList))  
    }  
  }  
}
```

Work on Streams - Qualifying

```
class QualifierFunction extends FlatMapFunction[FilteredEvent, QualifiedEvent] {  
  override def flatMap(value: FilteredEvent, out: Collector[QualifiedEvent]): Unit = {  
    Try(JsonPath.parse(value.event.payload)).map(ctx => {  
      value.controls.foreach(control => {  
        Try {  
          val result: String = ctx.read(control.jsonPath)  
          if (!result.isEmpty) {  
            out.collect(QualifiedEvent(value.event, control))  
          }  
        }  
      })  
    })  
  }  
}
```

Work on Streams - Counting

```
class CounterFunction extends FlatMapFunction[QualifiedEvent, ControlEvent] {  
  var counts = scala.collection.mutable.HashMap[String, Int]()  
  
  override def flatMap(value: QualifiedEvent, out: Collector[ControlEvent]): Unit = {  
    val key = s"${value.event.customerId}${value.control.alertId}"  
    if (counts.contains(key)) {  
      counts.put(key, counts.get(key).get + 1)  
      println(s"Count for ${key}: ${counts.get(key).get}")  
    } else {  
      counts.put(key, 1)  
      println(s"Count for ${key}: ${counts.get(key).get}")  
    }  
  }  
}
```

Wiring It Up

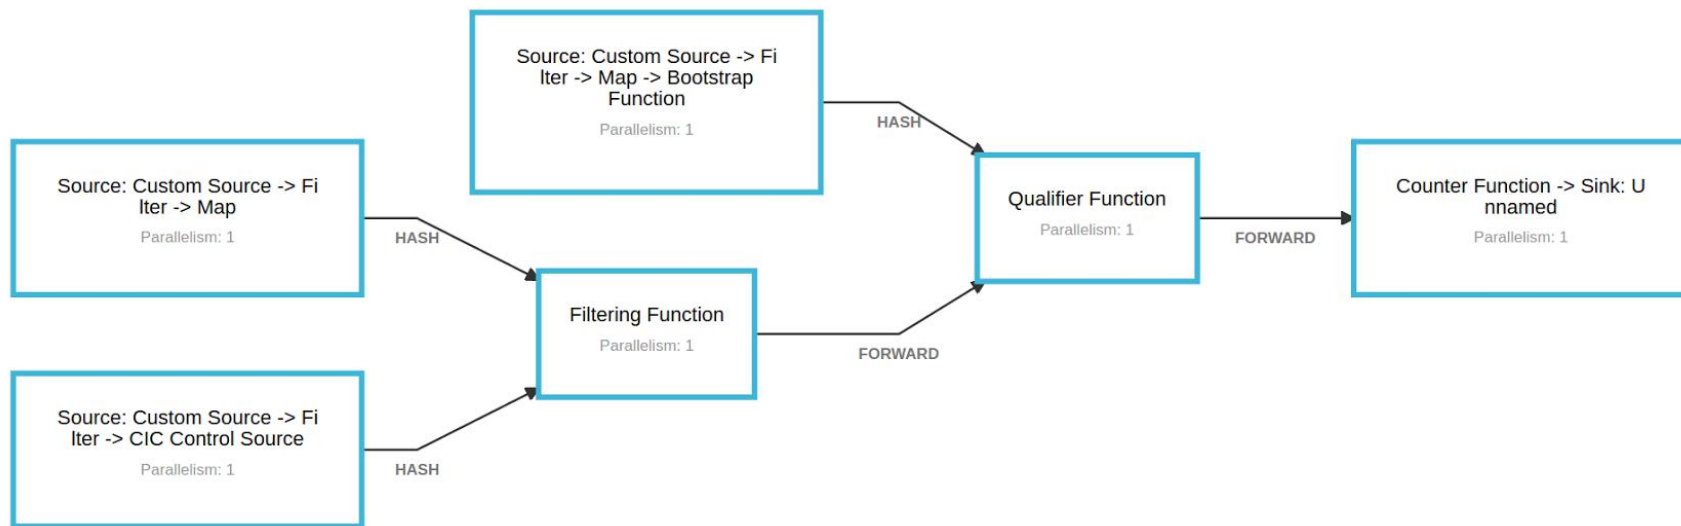
```
val filterFunction = new FilterFunction
val qualifierFunction = new QualifierFunction
val counterFunction = new CounterFunction

// Join the control and event streams
val filterStream = globalControlStream.union(specificControlStream)
    .connect(
        eventStream
    )
    .flatMap(filterFunction).name("Filtering Function")
    .flatMap(qualifierFunction).name("Qualifier Function")
    .flatMap(counterFunction).name("Counter Function")
```

What about historical data?

- The current solution works great against the live stream of data, but...
how do you get to the current number when the events we need have already gone through the system?

Basic Implementation



Counter Function Updated

```
override def flatMap(value: QualifiedEvent, out: Collector[ControlEvent]): Unit = {  
  val key = s"${value.event.customerId}${value.control.alertId}"  
  if (counts.contains(key)) {  
    counts.put(key, counts.get(key).get + 1)  
    println(s"Count for ${key}: ${counts.get(key).get}")  
  } else {  
    val c = value.control  
    counts.put(key, 1)  
    out.collect(ControlEvent(c.customerId, c.alertId, c.alertName, c.alertDescription,  
c.threshold, c.jsonPath, value.event.customerId))  
    println(s"Bootstrap count for ${key}: ${counts.get(key).get}")  
  }  
}
```



```
class BootstrapFunction extends FlatMapFunction[ControlEvent, FilteredEvent] {  
  override def flatMap(value: ControlEvent, out: Collector[FilteredEvent]): Unit = {  
    val stream = getClass.getResourceAsStream("/events.txt")  
  
    Source.fromInputStream(stream)  
      .getLines  
      .toList  
      .map(x => CustomerEvent(x))  
      .filter(x => x.customerId == value.bootstrapCustomerId)  
      .foreach(x => {  
        out.collect(FilteredEvent(x, List(value)))  
      })  
  }  
}
```

```
val bootstrapStream = env.addSource(new FlinkKafkaConsumer09("bootstrap", new ControlEventSchema(),  
properties))  
    .filter(x => x.isDefined).map(x => x.get).flatMap(bootstrapFunction).name("Bootstrap Function")  
    .keyBy((fe: FilteredEvent) => { fe.event.customerId } )
```

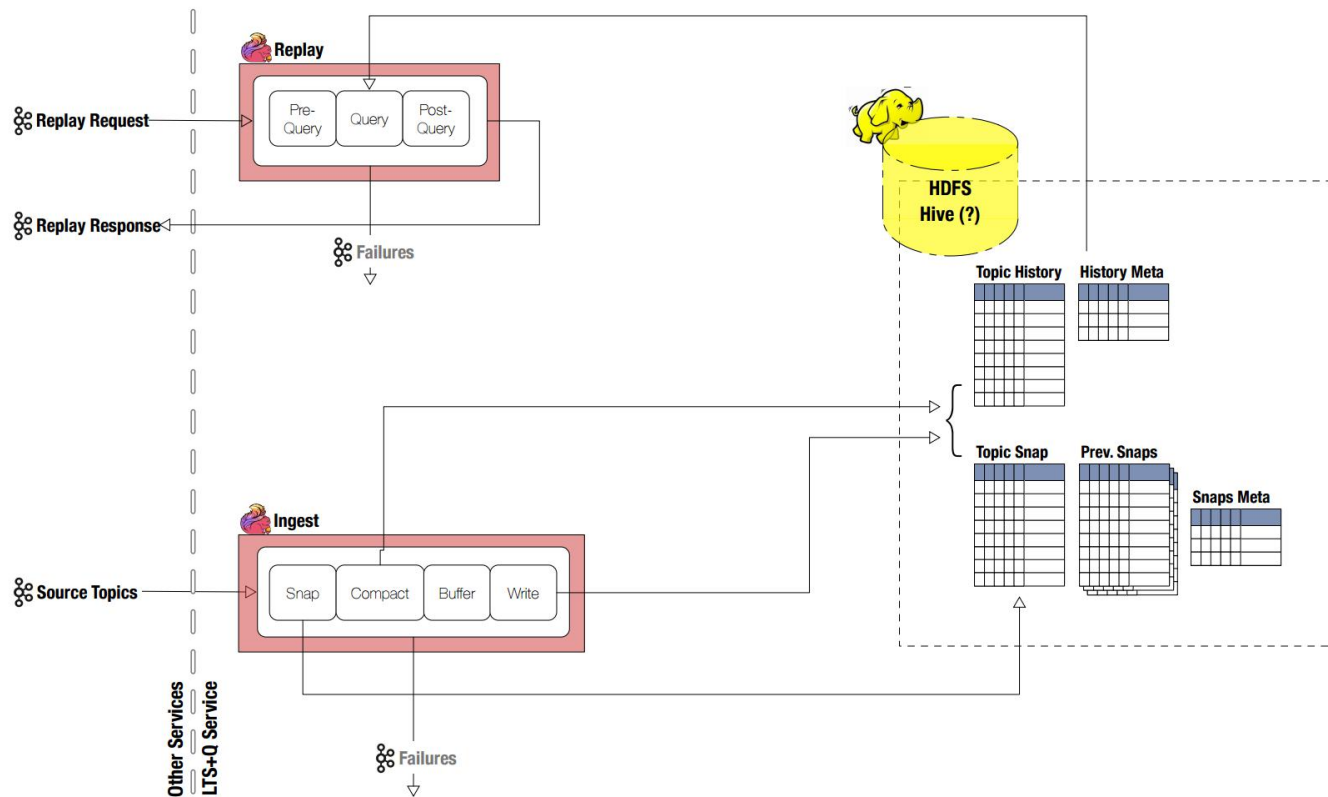
```
val bootstrapSink = new FlinkKafkaProducer09("bootstrap", new ControlEventSchema(), properties)
```

```
val filterStream = globalControlStream.union(specificControlStream)  
    .connect(eventStream)  
    .flatMap(filterFunction).name("Filtering Function")  
    .union(bootstrapStream)  
    .flatMap(qualifierFunction).name("Qualifier Function")  
    .flatMap(counterFunction).name("Counter Function")  
    .addSink(bootstrapSink)
```

- Extremely simplified compared to what you would want in production
 - A real system requires both ingestion of all events and some kind of query system (API, database, kafka, etc...)
 - Be careful with how you implement the query system. It is easy to block the whole stream with an API or database (Async I/O!!!!).
- There should be some kind of tracking on the event which triggered to bootstrap
 - You will need to get all events up until that event from the storage system
- If ordering does not matter:
 - Can just replay all previous events as in the example code

- If ordering does matter:
 - You will need to collect and block all incoming events for that specific key until the replay of old events completes

Our Solution For Historical Data - Long Term Storage



TO THE DEMO!



BetterCloud

Further Areas for Improvement

- Most expensive operations for us are JsonPath operations. It may be worth moving towards broadcasting all alerts instead of just global and move the keyed stream to as late as possible (we are switching to this)
- State is not being saved. At a minimum the controls and counts should be saved in state
- The current models only support incrementing
- JsonPath supports different types so you can have more operators than equals (>, <, etc...)
- Only a single property path is supported

We are hiring (go figure, right?)!!!

Seriously, we are always hiring. If you are interested talk to David Hardwick or visit <https://www.bettercloud.com/careers/>

Questions?

<https://github.com/brelloch/FlinkForward2017>

