



# No shard left behind

Dynamic Work Rebalancing  
and other adaptive features in  
Apache Beam

*Malo Denielou (malo@google.com)*

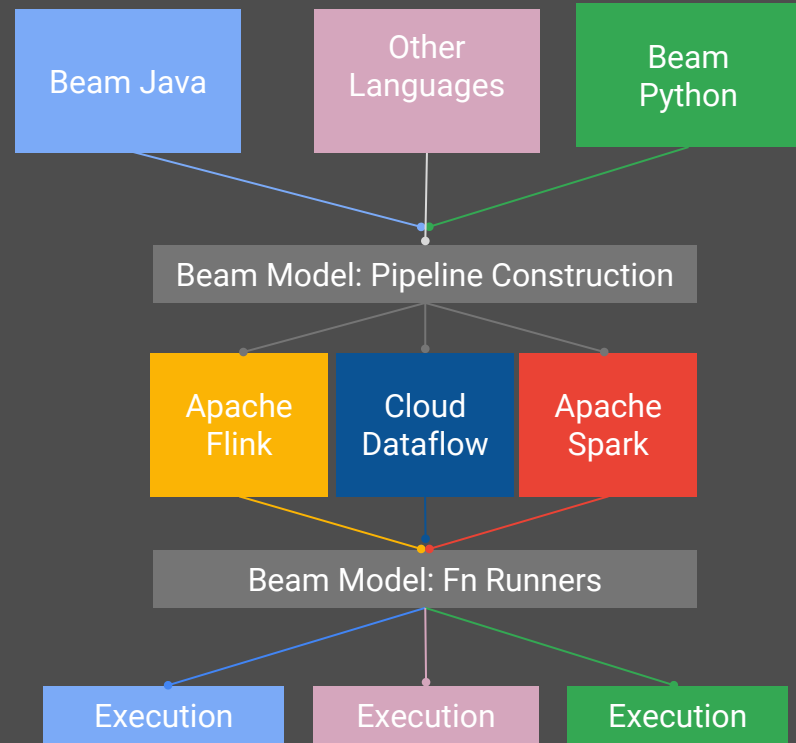





Apache Beam is a **unified** programming model designed to provide **efficient** and **portable** data processing pipelines.

# Apache Beam


1. The Beam Programming Model
2. SDKs for writing Beam pipelines -- Java/Python/...
3. Runners for existing distributed processing backends
  - Apache Flink
  - Apache Spark
  - Apache Apex
  - Dataflow
  - Direct runner (for testing)




# Apache Beam use cases




1. Classic Batch




2. Batch with Fixed Windows




3. Streaming



4. Streaming with Speculative + Late Data




5. Streaming With Retractions




6. Streaming With Sessions

# Data processing for realistic workloads




Streaming pipelines have variable input




Batch pipelines have stages of different sizes

# The curse of configuration




Over-provisioning resources?



Under-provisioning on purpose?


**A considerable effort is spent to finely tune all the parameters of the jobs.**

# Ideal case



**A system that adapts.**

# The straggler problem in batch



Tasks do not finish evenly on the workers.


- Data is not evenly distributed among tasks
- Processing time is uneven between tasks
- Runtime constraints

**Effects are cumulative per stage!**



# Common straggler mitigation techniques


- Split files into equal sizes?
- Pre-emptively over split?
- Detect slow workers and reexecute?
- Sample the data and split based on partial execution



**All have major costs, but do not solve completely the problem.**

# Common straggler mitigation techniques

- Split files into equal sizes?
- Pre-emptively over split?
- Detect slow workers and reexecute?
- Sample the data and split based on partial execution



**All have major costs, but do not solve completely the problem.**

« The most straightforward way to tune the number of partitions is experimentation: Look at the number of partitions in the parent RDD and then keep multiplying that by 1.5 until performance stops improving. »

From [blog]how-to-tune-your-apache-spark-jobs

**No amount of upfront heuristic tuning** (be it manual or automatic) is enough to guarantee good performance: the **system will always hit unpredictable situations** at run-time.

A system that's able to **dynamically adapt and get out of a bad situation** is much more powerful than one that **heuristically hopes to avoid** getting into it.

**Fine-tuning execution parameters** goes against having a truly **portable** and **unified** programming environment.



# Beam abstractions empower runners

A **bundle** is group of elements of a PCollection processed and committed together.

APIs (ParDo/DoFn):

- `setup()`
- `startBundle()`
- `processElement()` n times
- `finishBundle()`
- `teardown()`

**Streaming runner:**

- **small bundles**, low-latency **pipelining** across stages, **overhead** of frequent commits.

**Classic batch runner:**

- **large bundles**, fewer **large commits**, more **efficient**, **long synchronous stages**.

Other runner strategies may strike a different balance.

# Beam abstractions empower runners

## *Efficiency at runner's discretion*

“Read from this source, **splitting it 1000 ways**”

→ **user** decides

“Read from this source”

→ **runner** decides

APIs for portable Sources:

- `long getEstimatedSize()`
- `List<Source> splitIntoBundles(size)`

# Beam abstractions empower runners

## *Efficiency at runner's discretion*

“Read from this source, **splitting it 1000 ways**”


→ **user** decides

“Read from this source”

→ **runner** decides

APIs:

- `long getEstimatedSize()`
- `List<Source> splitIntoBundles(size)`



# Beam abstractions empower runners

## *Efficiency at runner's discretion*

“Read from this source, **splitting it 1000 ways**”


→ **user** decides

“Read from this source”

→ **runner** decides

APIs:

- `long getEstimatedSize()`
- `List<Source> splitIntoBundles(size)`



# Beam abstractions empower runners

## *Efficiency at runner's discretion*

“Read from this source, **splitting it 1000 ways**”

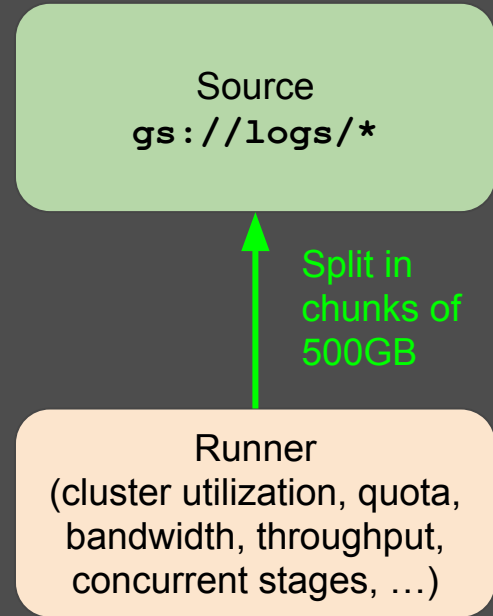
→ **user** decides

“Read from this source”

→ **runner** decides

APIs:

- `long getEstimatedSize()`
- `List<Source> splitIntoBundles(size)`





# Beam abstractions empower runners

## *Efficiency at runner's discretion*

"Read from this source, **splitting it 1000 ways**"

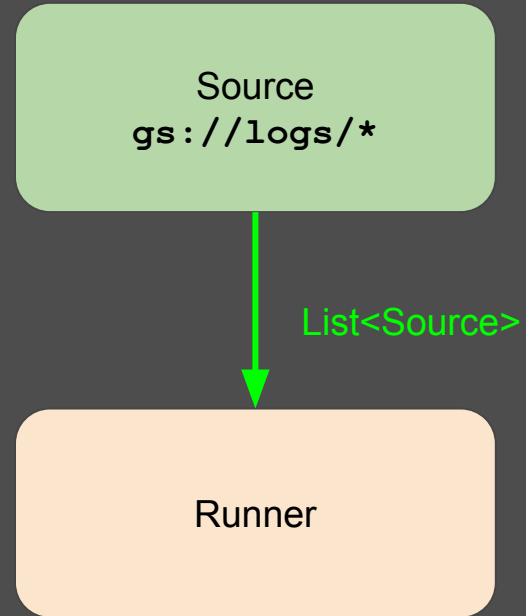
→ **user** decides

"Read from this source"

→ **runner** decides

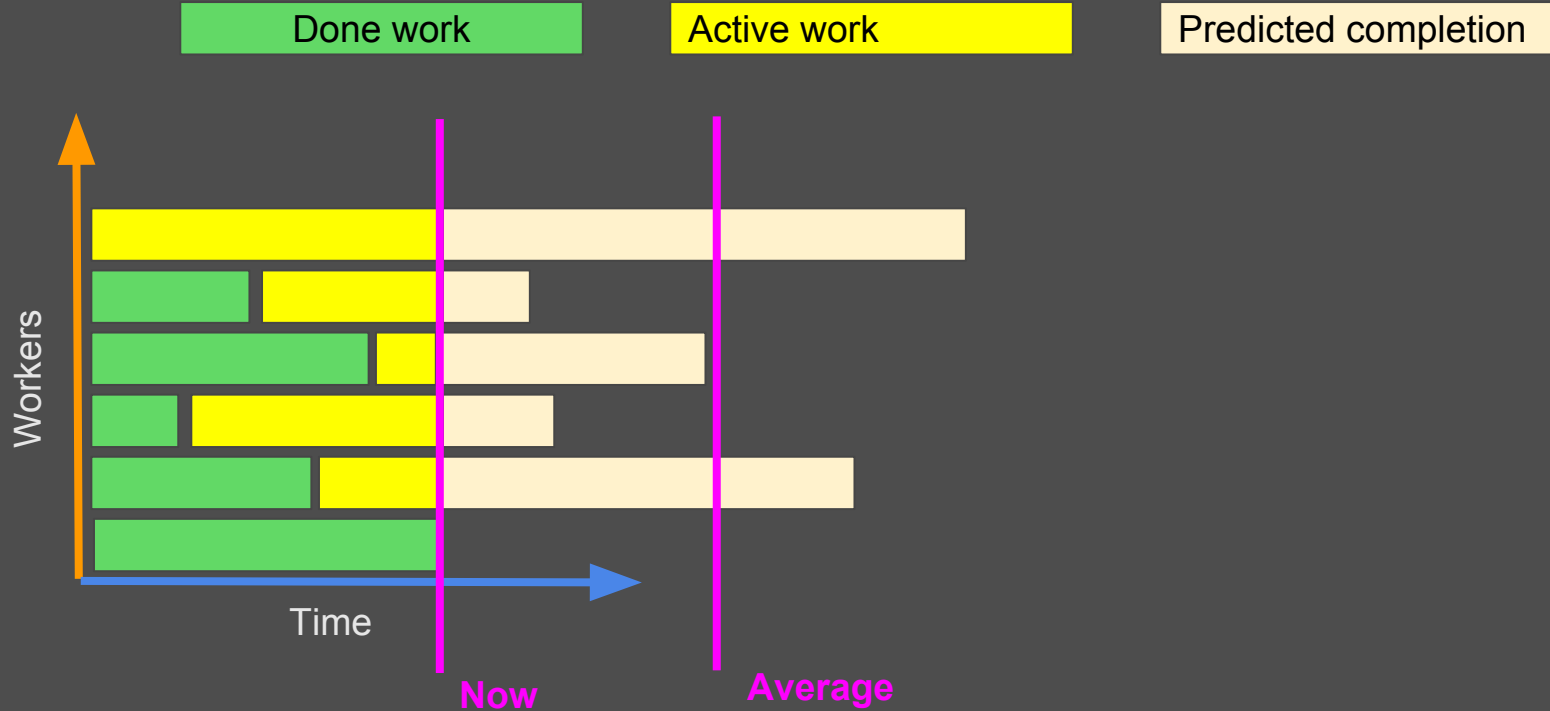
APIs:

- `long getEstimatedSize()`
- `List<Source> splitIntoBundles(size)`

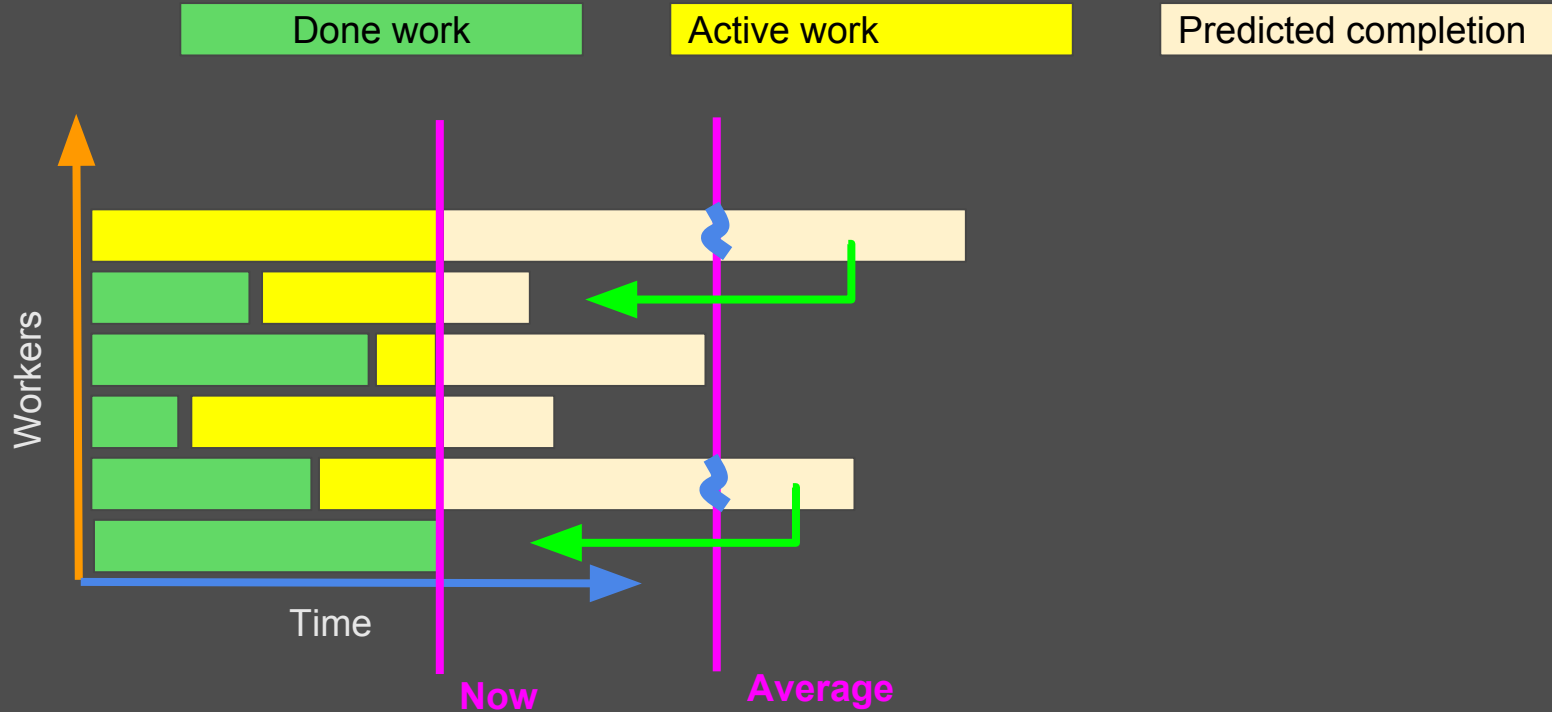


# Solving the straggler problem: Dynamic Work Rebalancing

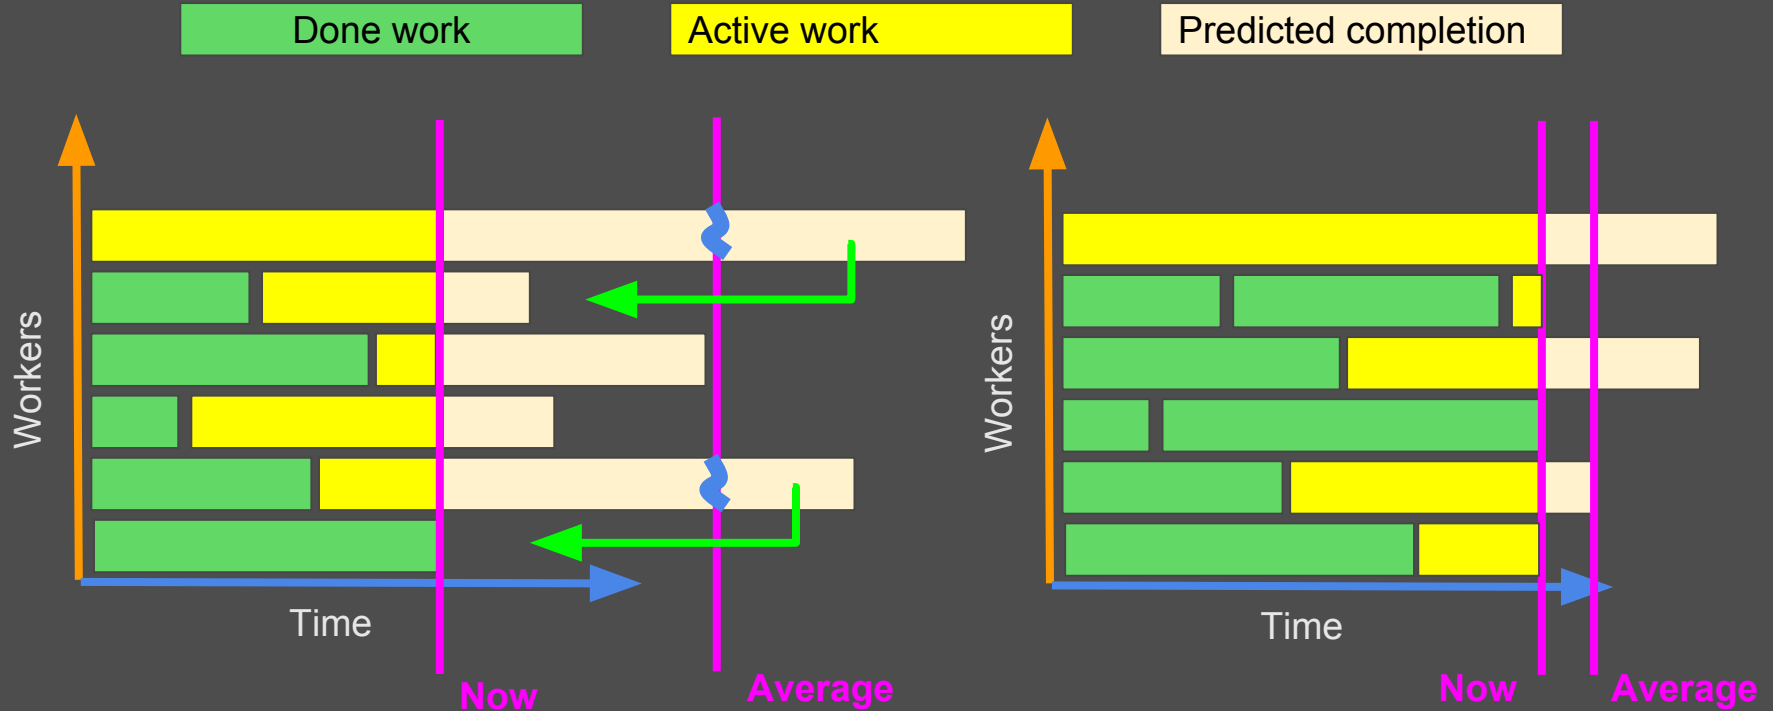
# Solving the straggler problem: Dynamic Work Rebalancing



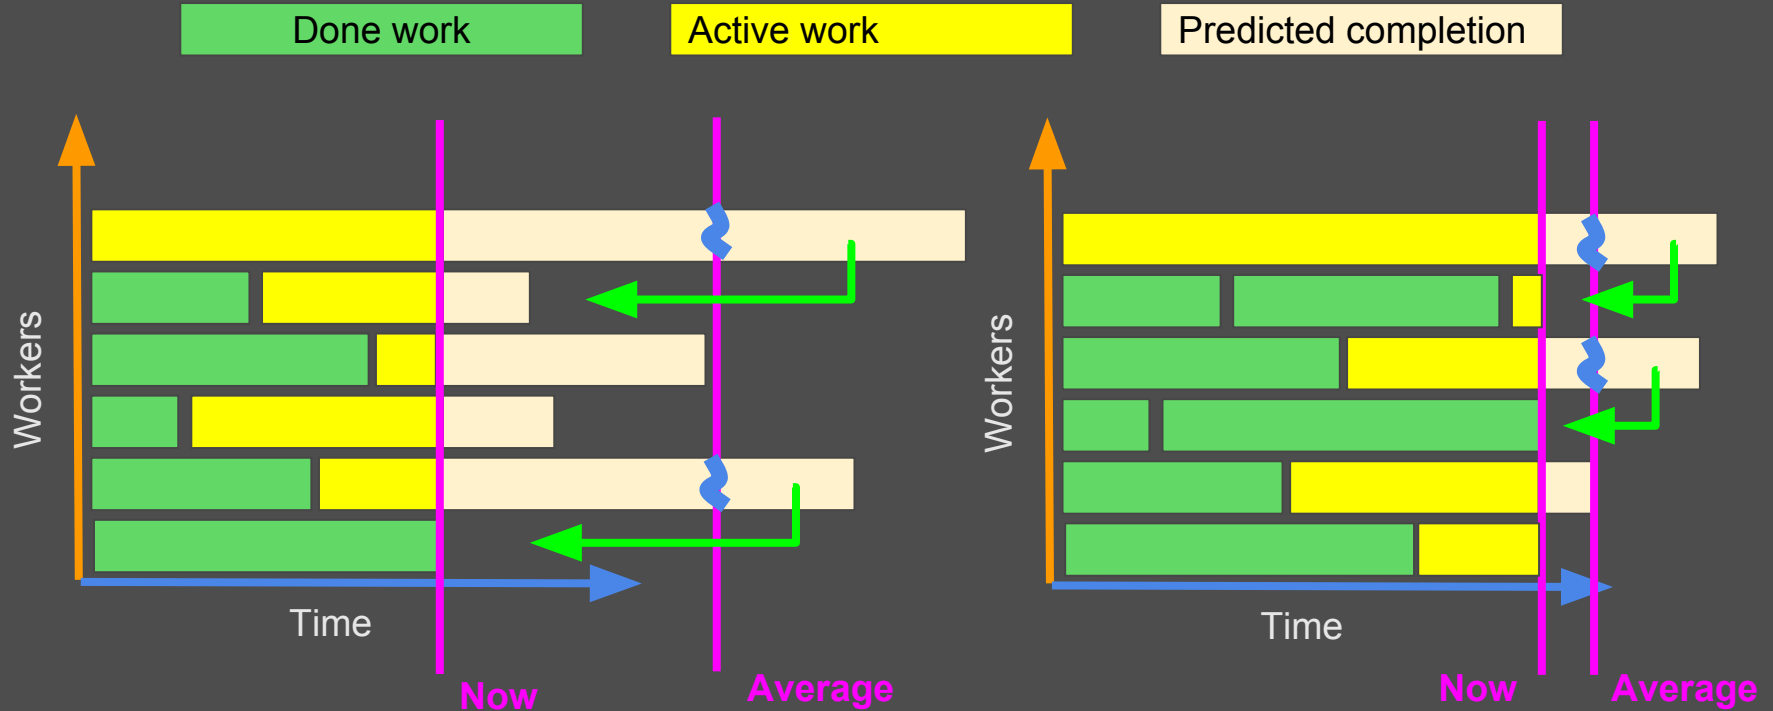
# Solving the straggler problem: Dynamic Work Rebalancing



# Solving the straggler problem: Dynamic Work Rebalancing



# Solving the straggler problem: Dynamic Work Rebalancing



# Dynamic Work Rebalancing in the wild



*A classic MapReduce job (read from Google Cloud Storage, GroupByKey, write to Google Cloud Storage), 400 workers.*

*Dynamic Work Rebalancing disabled to demonstrate stragglers.*


*X axis: time (total ~20min.); Y axis: workers*

# Dynamic Work Rebalancing in the wild



*A classic MapReduce job (read from Google Cloud Storage, GroupByKey, write to Google Cloud Storage), 400 workers. Dynamic Work Rebalancing disabled to demonstrate stragglers.*

*X axis: time (total ~20min.); Y axis: workers*



*Same job,  
Dynamic Work Rebalancing enabled.  
X axis: time (total ~15min.); Y axis: workers*

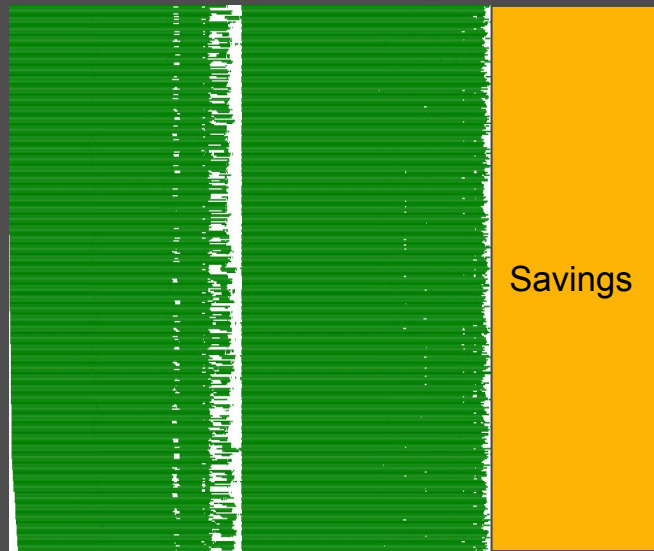


# Dynamic Work Rebalancing in the wild



*A classic MapReduce job (read from Google Cloud Storage, GroupByKey, write to Google Cloud Storage), 400 workers. Dynamic Work Rebalancing disabled to demonstrate stragglers.*

*X axis: time (total ~20min.); Y axis: workers*



*Same job,  
Dynamic Work Rebalancing enabled.  
X axis: time (total ~15min.); Y axis: workers*


# Dynamic Work Rebalancing with Autoscaling

**Initial allocation** of 80 workers, based on **size**

Multiple rounds of **upsizing**, enabled by dynamic work rebalancing

**Upscales to 1000 workers.**

- tasks are balanced
- no oversplitting or manual tuning



# Apache Beam enable dynamic adaptation

Beam Source Readers provide **simple progress signals**, which enable runners to take action based on execution-time characteristics.

**All Beam runners can implement Autoscaling and Dynamic Work Rebalancing.**

APIs for how much work is pending.

- **bounded:** `double getFractionConsumed()`
- **unbounded:** `long getBacklogBytes()`

APIs for splitting:

- **bounded:**
  - `Source splitAtFraction(double)`
  - `int getParallelismRemaining()`
- **unbounded:**
  - Coming soon ...



Apache Beam is a **unified** programming model designed to provide **efficient** and **portable** data processing pipelines.

## To learn more

Read our blog posts!

- No shard left behind: Dynamic work rebalancing in Google Cloud Dataflow  
<https://cloud.google.com/blog/big-data/2016/05/no-shard-left-behind-dynamic-work-rebalancing-in-google-cloud-dataflow>
- Comparing Cloud Dataflow autoscaling to Spark and Hadoop  
<https://cloud.google.com/blog/big-data/2016/03/comparing-cloud-dataflow-autoscaling-to-spark-and-hadoop>

Join the Apache Beam community!

<https://beam.apache.org/>

